



國立交通大學
National Chiao Tung University

Kernel, Gaussian Process, SVM

ML 2021 Fall

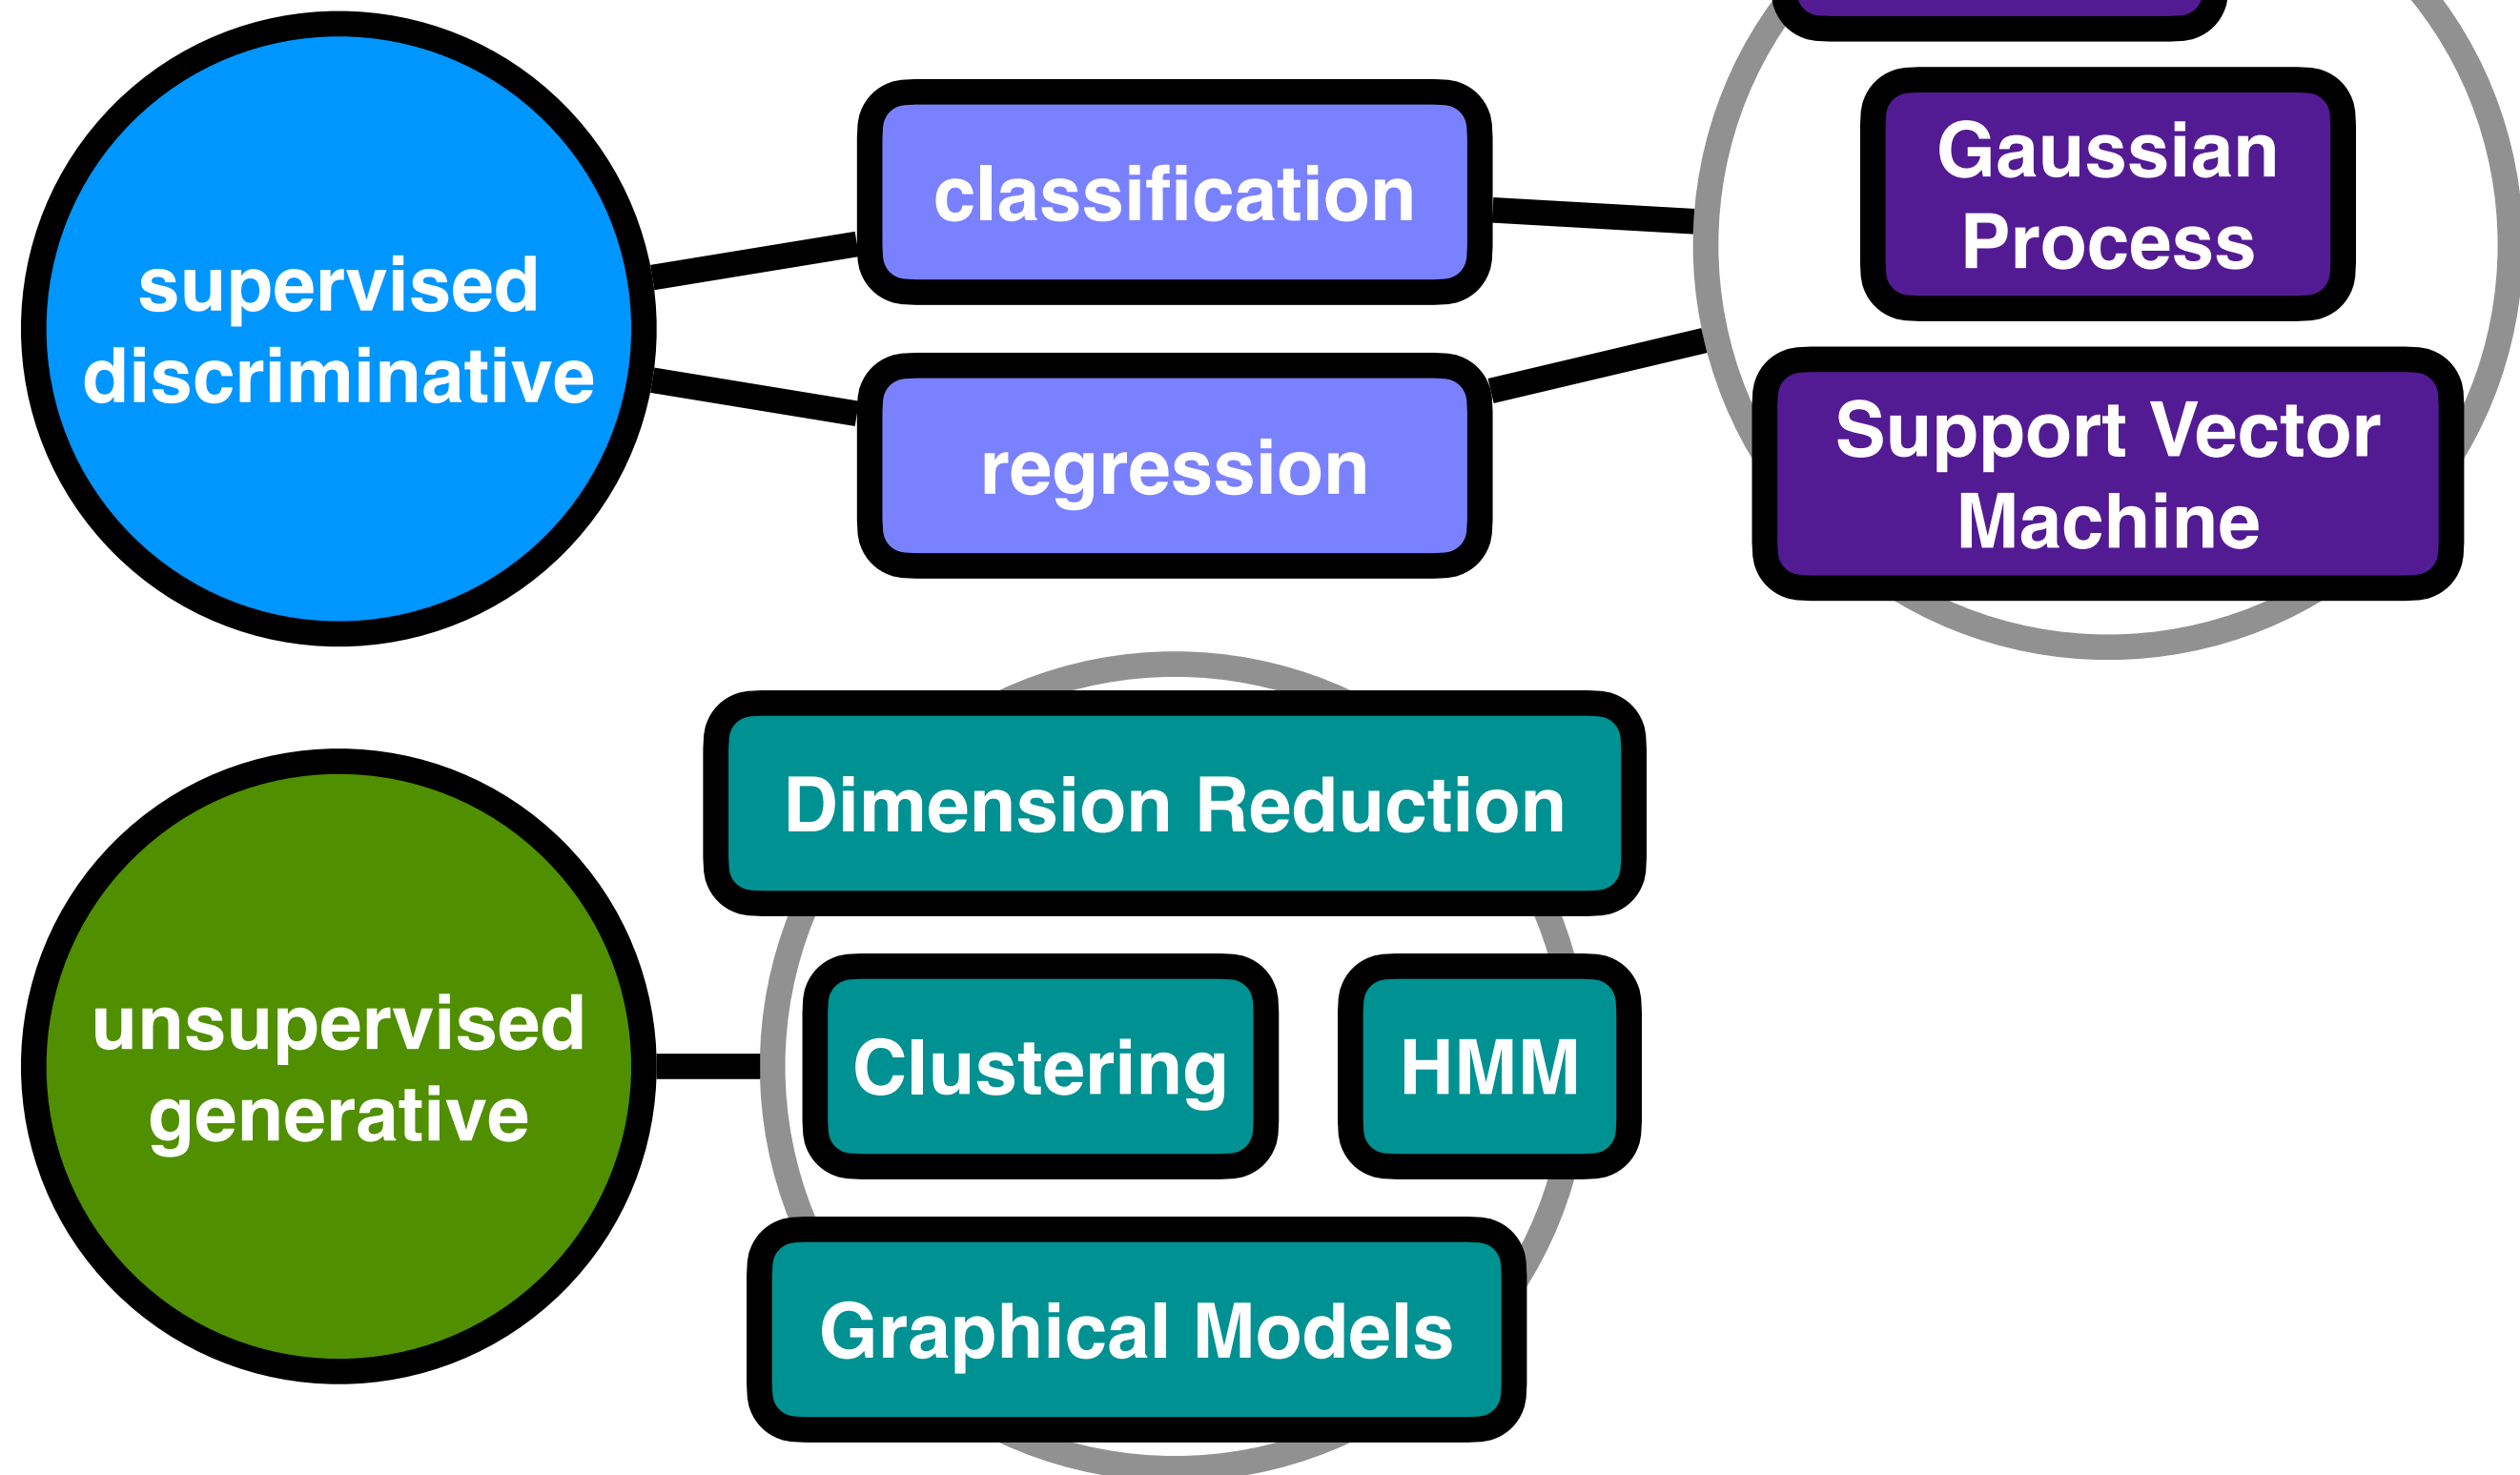
Wei-Chen (Walon) Chiu

邱維辰

Enriched Vision Applications
Laboratory

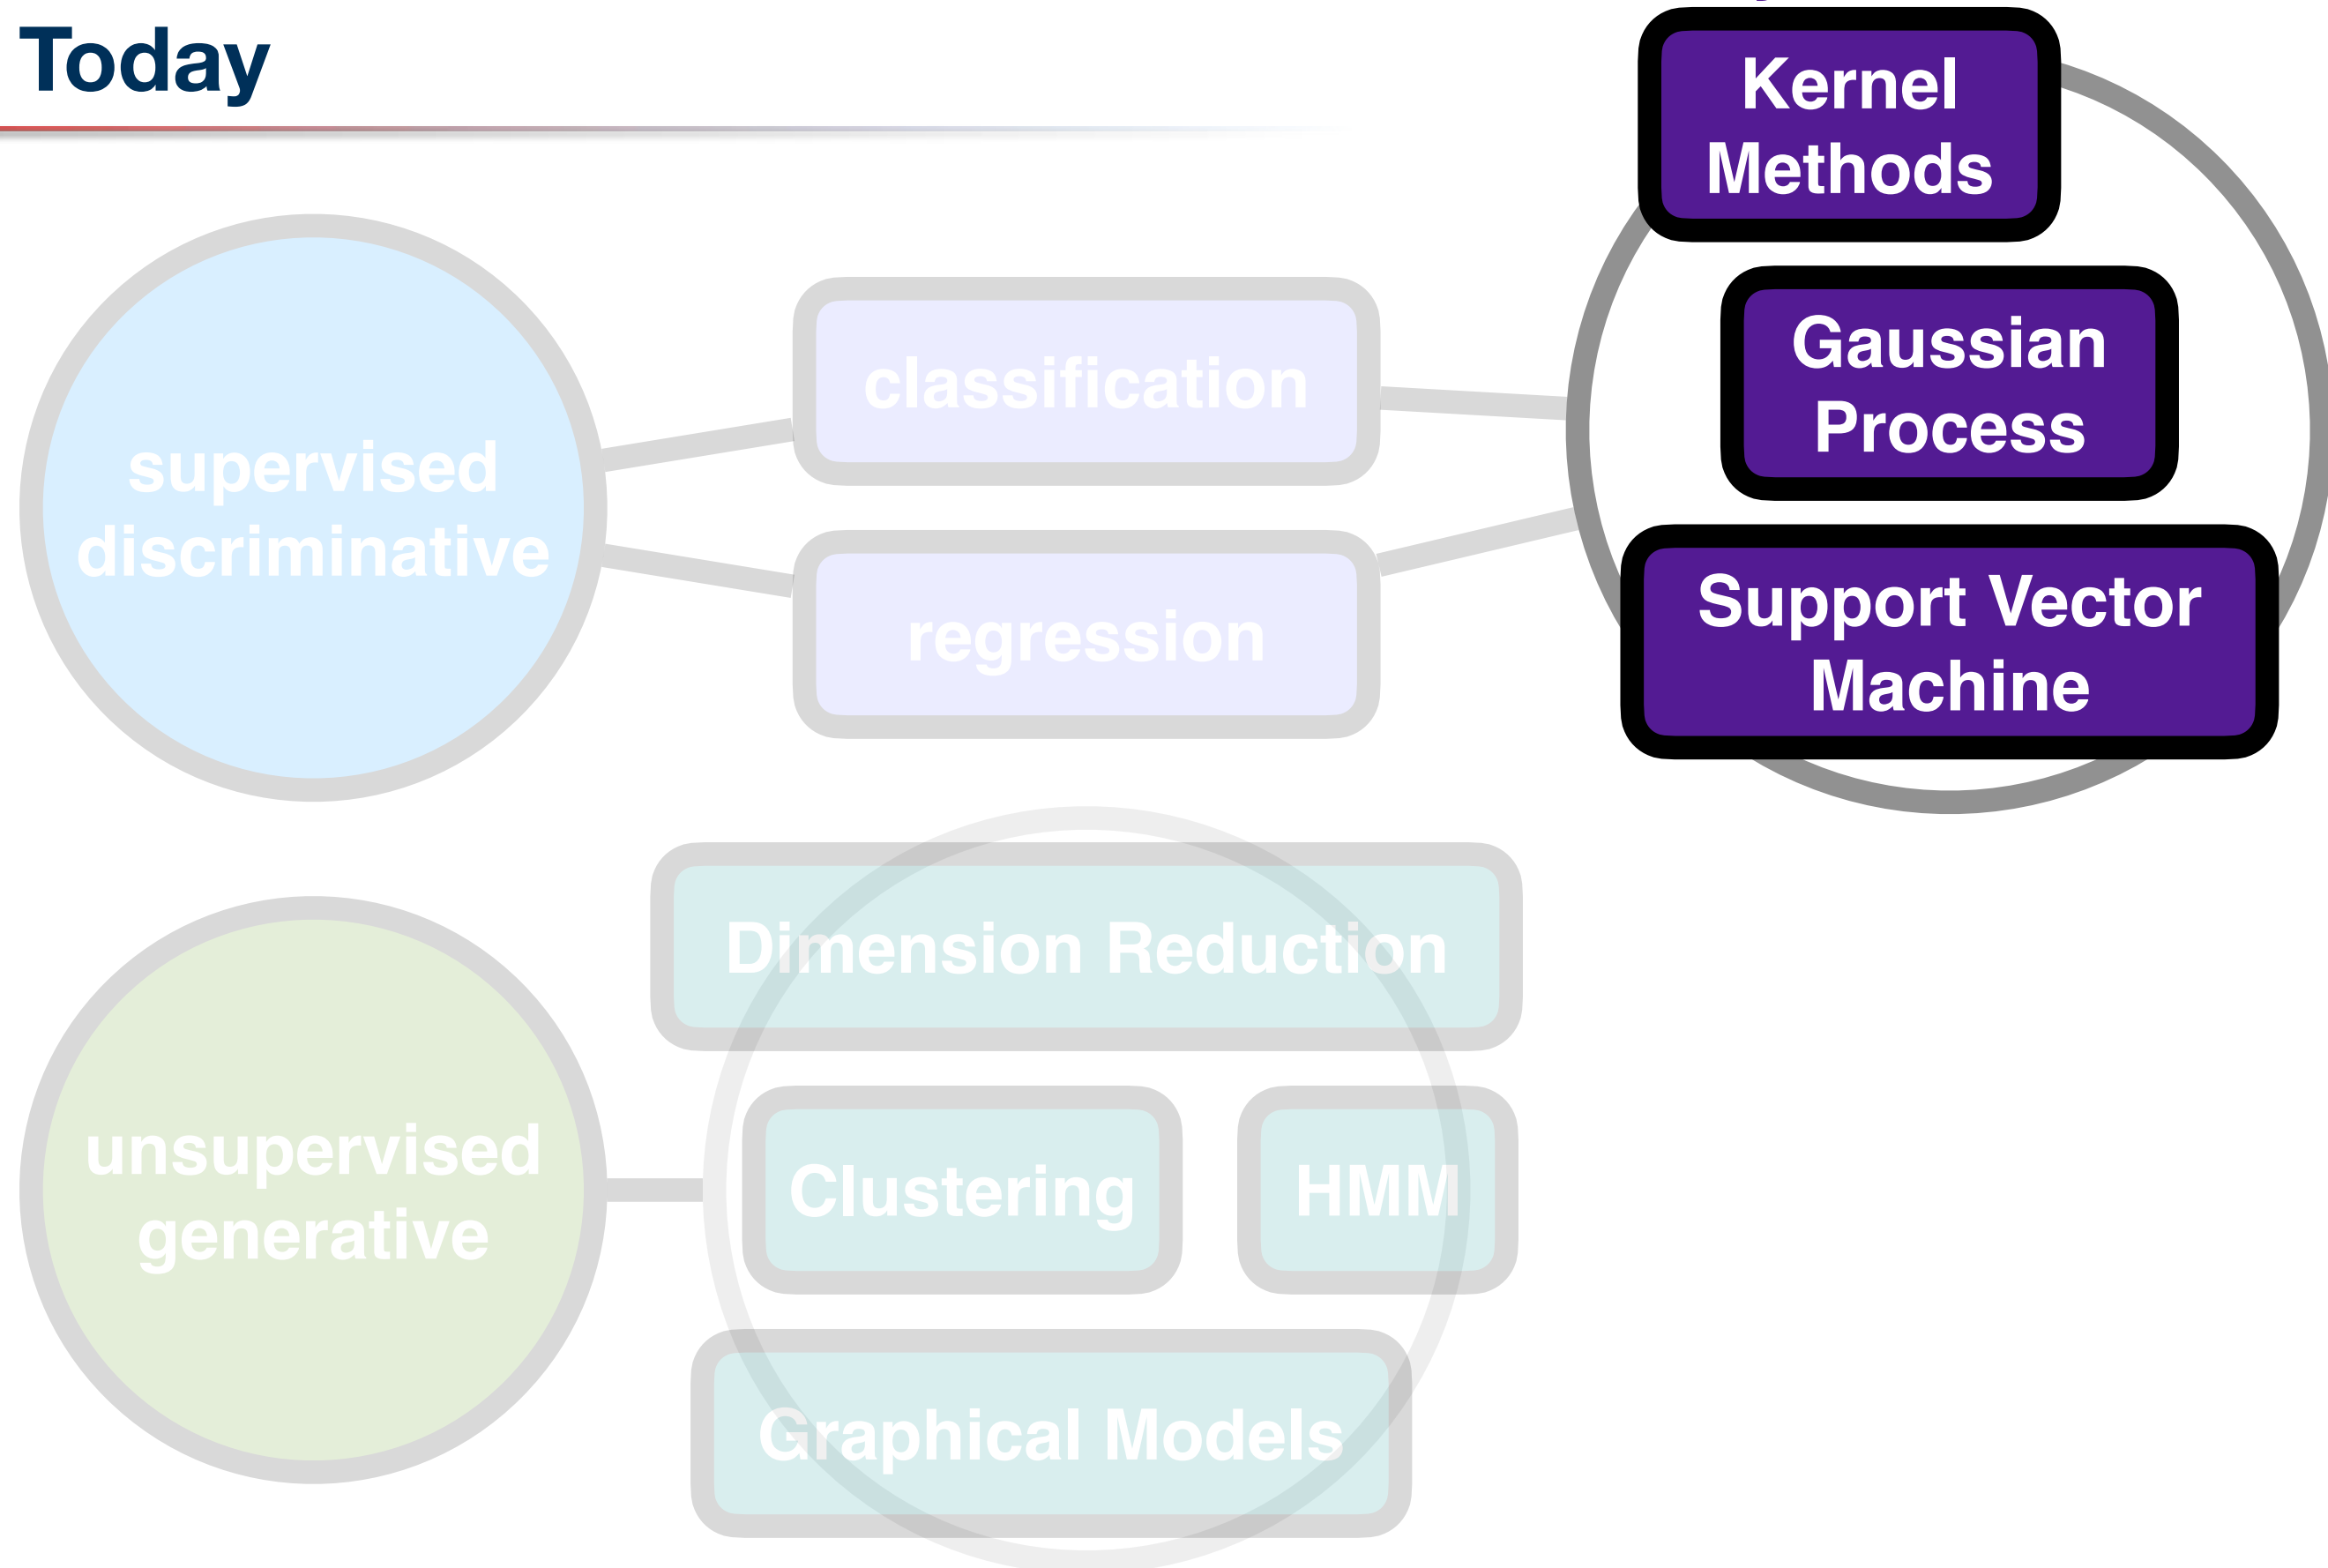


Topics to be tackled around the corner

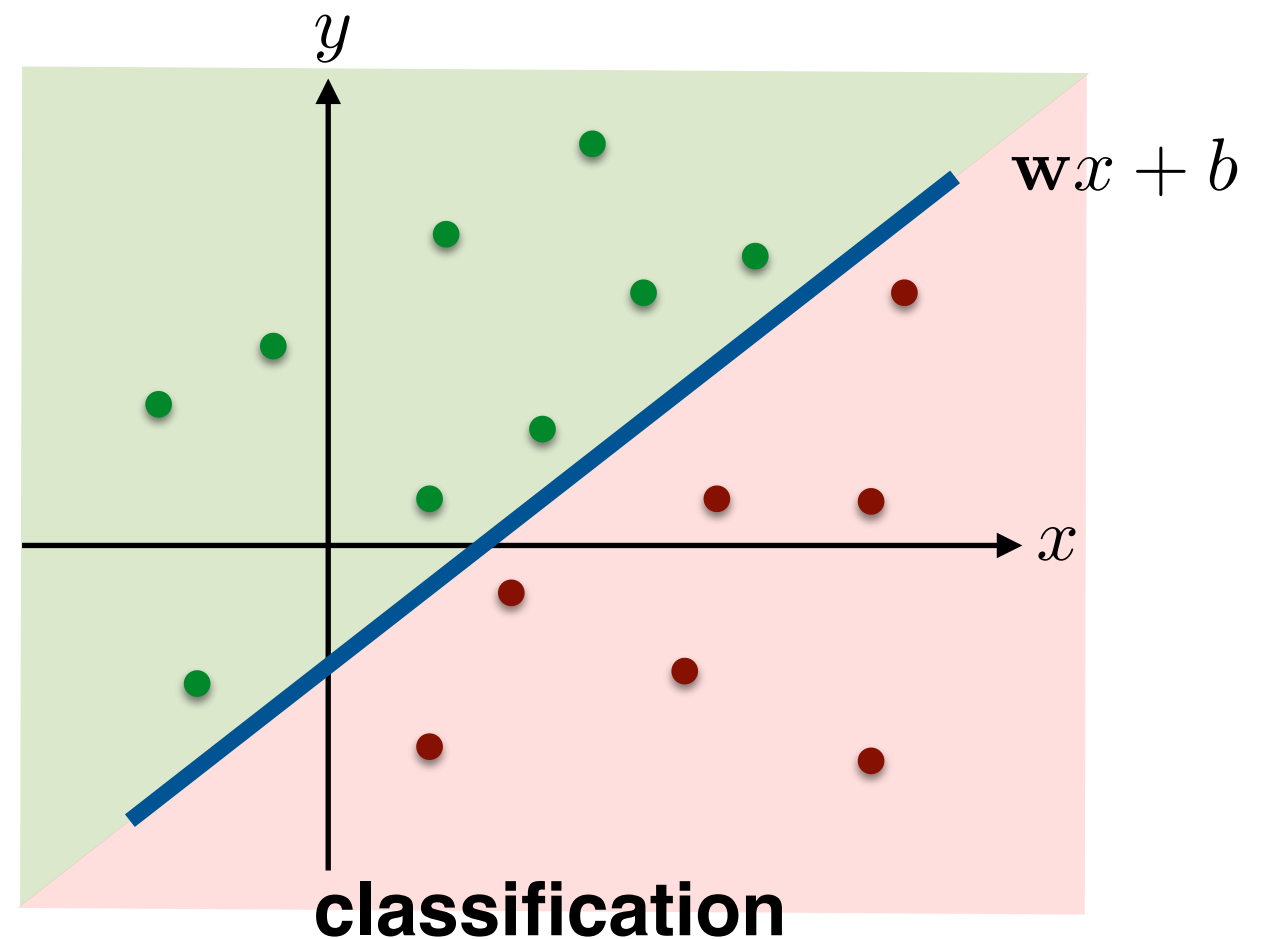



Today

memory-based methods



What you guys have learnt




- No matter linear parametric models for regression or classification, throw away the training data after learning the parameters of the model.
- Make prediction based on those parameters only.

Another way of thinking


- Training data is valuable, try to keep it even till prediction
 - **Memory-based methods**
 - **We would like to utilise the training data for prediction!**
- For instance, KNN, Parzen

Another way of thinking

- Training data is valuable, try to keep it even till prediction
 - **Memory-based methods**
 - **We would like to utilise the training data for prediction!**
- For instance, **KNN**, Parzen probability estimation




dataset




NN (nearest neighbour):
find the closest training data point,
take its class as prediction

Another way of thinking


- Training data is valuable, try to keep it even till prediction
 - **Memory-based methods**
 - **We would like to utilise the training data for prediction!**
- For instance, **KNN**, Parzen probability estimation



dataset




NN (nearest neighbour):




KNN:
prediction based on
K closest training data
points, do majority vote

Another way of thinking

- Training data is valuable, try to keep it even till prediction
 - **Memory-based methods**
 - **We would like to utilise the training data for prediction!**
- For instance, KNN, **Parzen probability estimation**




dataset




put  on each training data points

Another way of thinking

- Training data is valuable, try to keep it even till prediction
 - **Memory-based methods**
 - **We would like to utilise the training data for prediction!**
 - **Need to evaluate the distance from “testing” to “training” data**




KNN



Parzen probability estimation

Data in the feature space

- General machine learning scheme



- ▶ project x to a feature space by feature mapping ϕ
- ▶ we need to evaluate the distance/similarity between data $\phi(x)$ in the feature space 🖐 **dot product** is somehow related to similarity

$$\phi(\mathbf{x})^\top \phi(\mathbf{x}')$$

- ▶ **kernel** function!


$$\phi(\mathbf{x})^\top \phi(\mathbf{x}') = k(\mathbf{x}, \mathbf{x}')$$

- ▶ ... so what? why do we need this?

Example for kernel to exist

- Linear regression with L_2 regularisation


$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$



Example for kernel to exist

- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$




- with $\frac{\partial J}{\partial \mathbf{w}} = 0$, we can see that \mathbf{w} is the linear combination of $\phi(\mathbf{x})$

$$\mathbf{w} = -\frac{1}{\lambda} \sum_{n=1}^N \underbrace{\{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}}_{\text{coefficients}} \phi(\mathbf{x}_n) = \sum_{n=1}^N \underbrace{a_n}_{\text{coefficients}} \phi(\mathbf{x}_n) = \Phi^T \mathbf{a}$$

Example for kernel to exist

- Linear regression with L_2 regularisation



$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- with $\frac{\partial J}{\partial \mathbf{w}} = 0$, we can see that \mathbf{w} is the linear combination of $\phi(\mathbf{x})$

$$\mathbf{w} = -\frac{1}{\lambda} \sum_{n=1}^N \underbrace{\{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}}_{\text{coefficients}} \phi(\mathbf{x}_n) = \sum_{n=1}^N \underbrace{a_n}_{\text{coefficients}} \phi(\mathbf{x}_n) = \Phi^T \mathbf{a}$$

- put $\mathbf{w} = \Phi^T \mathbf{a}$ back to J

$$J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \underbrace{\Phi \Phi^T \Phi \Phi^T}_{\text{kernel}} \mathbf{a} - \mathbf{a}^T \underbrace{\Phi \Phi^T}_{\text{kernel}} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \underbrace{\Phi \Phi^T}_{\text{kernel}} \mathbf{a}$$

helloworld! kernel $\phi(\mathbf{x})^T \phi(\mathbf{x}') = k(\mathbf{x}, \mathbf{x}')$

Example for kernel to exist

- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$


here comes dual form!

- put $\mathbf{w} = \Phi^T \mathbf{a}$ back to J

$$J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \underline{\Phi \Phi^T \Phi \Phi^T} \mathbf{a} - \mathbf{a}^T \underline{\Phi \Phi^T} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \underline{\Phi \Phi^T} \mathbf{a}$$

- $K_{nm} = \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m)$ Gram matrix: $K = \Phi \Phi^T$

$$J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \mathbf{K} \mathbf{K} \mathbf{a} - \mathbf{a}^T \mathbf{K} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \mathbf{K} \mathbf{a}.$$



Example for kernel to exist

- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

dual form


▶ $J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \mathbf{K} \mathbf{K} \mathbf{a} - \mathbf{a}^T \mathbf{K} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \mathbf{K} \mathbf{a}.$

▶ with $\frac{\partial J}{\partial \mathbf{a}} = 0$, get $\mathbf{a} = (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}.$

- ▶ when performing prediction:

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) = \mathbf{a}^T \Phi \phi(\mathbf{x}) = \mathbf{k}(\mathbf{x})^T (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}$$

where vector $\mathbf{k}(\mathbf{x})$ with elements $k_n(\mathbf{x}_n) = k(\mathbf{x}_n, \mathbf{x})$



Example for kernel to exist

- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

dual form


▶ $J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \mathbf{K} \mathbf{K} \mathbf{a} - \mathbf{a}^T \mathbf{K} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \mathbf{K} \mathbf{a}.$

▶ with $\frac{\partial J}{\partial \mathbf{a}} = 0$, get $\mathbf{a} = (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}.$

- ▶ when performing prediction:

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) = \mathbf{a}^T \Phi \phi(\mathbf{x}) = \mathbf{k}(\mathbf{x})^T (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}$$

where vector $\mathbf{k}(\mathbf{x})$ with elements $k_n(\mathbf{x}_n) = k(\mathbf{x}_n, \mathbf{x})$



**prediction based on
linear combination of
kernel functions evaluated
at training data points**


Example for kernel to exist

- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

dual form

$$J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \mathbf{K} \mathbf{K} \mathbf{a} - \mathbf{a}^T \mathbf{K} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \mathbf{K} \mathbf{a}.$$




ah-ha!

- with $\frac{\partial J}{\partial \mathbf{a}} = 0$, get $\mathbf{a} = (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}$.
- if \mathbf{x} always appears as scalar product, it might be good to try kernel!
known as “kernel trick”
a.k.a. “kernel substitution”*


Feature, feature

- All about features:
 - **projection from original data space to feature space**
 - How to choose a proper feature space?
 - for instance: **easier separation between classes** (classification)
 - prior knowledge about the class of functions to be learned
 - subsets of a basis of the function space (e.g. Fourier, Wavelet, etc.)




Feature, feature

- All about features:
 - **projection from original data space to feature space**
 - How to choose a proper feature space?
 - for instance: easier separation between classes (classification)
 - prior knowledge about the class of functions to be learned
 - subsets of a basis of the function space (e.g. Fourier, Wavelet, etc.)




- Sometimes we don't know what is the proper feature mapping...
 - still want to use memory-based methods (training data in prediction)
 - perhaps just some kernel? don't care what exactly is the feature map?



Some examples of kernel functions

- a very simple form $k(x, x') = e^{xx'}$




Some examples of kernel functions

- a very simple form $k(x, x') = e^{xx'}$
 - the corresponding feature map is

$$\forall x \in \mathbb{R}, r \in \mathbb{N} : \phi_r(x) = \frac{1}{\sqrt{r!}} x^r$$

$$k(x, x') = \sum_{r=0}^{\infty} \phi_r(x) \phi_r(x') = \sum_{r=0}^{\infty} \frac{x^r}{\sqrt{r!}} \frac{(x')^r}{\sqrt{r!}} = \sum_{r=0}^{\infty} \frac{(xx')^r}{r!} = e^{xx'}$$

simple kernel but it has even countably infinitely many feature maps!



Some examples of kernel functions

- a very simple form $k(x, x') = e^{xx'}$
 - the corresponding feature map is


$$\forall x \in \mathbb{R}, r \in \mathbb{N} : \phi_r(x) = \frac{1}{\sqrt{r!}} x^r$$

$$k(x, x') = \sum_{r=0}^{\infty} \phi_r(x) \phi_r(x') = \sum_{r=0}^{\infty} \frac{x^r}{\sqrt{r!}} \frac{(x')^r}{\sqrt{r!}} = \sum_{r=0}^{\infty} \frac{(xx')^r}{r!} = e^{xx'}$$

simple kernel but it has even countably infinitely many feature maps!



PS: simplest kernel is identity mapping
 $\phi(x) = x, \quad k(x, x') = x^\top x'$
it is called linear kernel.




Some examples of kernel functions

- Given $x_i \in \mathbb{R}^3, \phi(x_i) \in \mathbb{R}^{10}$

$$\phi(x_i) = \left[1, \sqrt{2}(x_i)_1, \sqrt{2}(x_i)_2, \sqrt{2}(x_i)_3, (x_i)_1^2, (x_i)_2^2, (x_i)_3^2, \sqrt{2}(x_i)_1(x_i)_2, \sqrt{2}(x_i)_1(x_i)_3, \sqrt{2}(x_i)_2(x_i)_3 \right]^\top$$

- kernel function: $\phi(x_i)^\top \phi(x_j) = \underline{(1 + x_i^\top x_j)^2}$

feature map projects 3-D to 10-D, but it results as a very simple form to compute kernel!



Some examples of kernel functions

- Radial Basis Kernel $e^{-\gamma \|x_i - x_j\|}$

▸ assume $x \in \mathbb{R}^1$ and $\gamma > 0$


$$e^{-\gamma \|x_i - x_j\|} = e^{-\gamma (x_i - x_j)^2} = e^{-\gamma x_i^2 + 2\gamma x_i x_j - \gamma x_j^2}$$

$$= e^{-\gamma x_i^2 - \gamma x_j^2} \left(1 + \frac{2\gamma x_i x_j}{1!} + \frac{(2\gamma x_i x_j)^2}{2!} + \frac{(2\gamma x_i x_j)^3}{3!} + \dots \right) \text{power series definition}$$

$$= e^{-\gamma x_i^2 - \gamma x_j^2} \left(1 \cdot 1 + \sqrt{\frac{2\gamma}{1!}} x_i \times \sqrt{\frac{2\gamma}{1!}} x_j + \sqrt{\frac{(2\gamma)^2}{2!}} x_i^2 \times \sqrt{\frac{(2\gamma)^2}{2!}} x_j^2 + \dots \right)$$

$$= \phi(x_i)^\top \phi(x_j)$$

where $\phi(x) = e^{-\gamma x^2} \left[1, \sqrt{\frac{2\gamma}{1!}} x, \sqrt{\frac{(2\gamma)^2}{2!}} x^2, \sqrt{\frac{(2\gamma)^3}{3!}} x^3, \dots \right]^\top$ **simple kernel but infinite-D feature map**



Example for kernel to exist


- Linear regression with L_2 regularisation

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ \mathbf{w}^T \phi(\mathbf{x}_n) - t_n \}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

↕
dual form
↕

▸ $J(\mathbf{a}) = \frac{1}{2} \mathbf{a}^T \mathbf{K} \mathbf{a} - \mathbf{a}^T \mathbf{K} \mathbf{t} + \frac{1}{2} \mathbf{t}^T \mathbf{t} + \frac{\lambda}{2} \mathbf{a}^T \mathbf{K} \mathbf{a}.$



▸ with $\frac{\partial J}{\partial \mathbf{a}} = 0$, get $\mathbf{a} = (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{t}.$



ah-ha!

*even it seems more complicated to
do matrix inversion in higher-D space,
but we can directly work on kernels (usually simple)
and avoid explicit introduction of feature vector $\phi(\mathbf{x})$,
which allows us implicitly to use **feature spaces**
of high, even infinite, dimensionality!*

Why High Dimensional Feature Space



How to construct valid kernel?

- a necessary and sufficient condition:
the **Gram matrix**: $K = \Phi\Phi^\top$ whose element are given by
 $k(x, x') = \phi(x)^\top \phi(x')$ is positive semidefinite

A symmetric function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is called a **positive ^{semi}definite kernel** if, for any finite J , any $x \in \mathcal{X}^J$ and any $c \in \mathbb{R}^J$:

$$0 \leq \sum_{i \in J} \sum_{j \in J} c_i c_j k(x_i, x_j)$$

The set of all real-valued positive definite kernels on \mathcal{X} is denoted $\mathbb{R}_+^{\mathcal{X} \times \mathcal{X}}$

How to construct valid kernel?

- a necessary and sufficient condition:
the **Gram matrix**: $K = \Phi\Phi^\top$ whose element are given by $k(x, x') = \phi(x)^\top \phi(x')$ is positive semidefinite
- You can build kernel from kernel
$$\begin{aligned}k(\mathbf{x}, \mathbf{x}') &= ck_1(\mathbf{x}, \mathbf{x}') \\k(\mathbf{x}, \mathbf{x}') &= f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}') \\k(\mathbf{x}, \mathbf{x}') &= q(k_1(\mathbf{x}, \mathbf{x}')) \\k(\mathbf{x}, \mathbf{x}') &= \exp(k_1(\mathbf{x}, \mathbf{x}')) \\k(\mathbf{x}, \mathbf{x}') &= k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}') \\k(\mathbf{x}, \mathbf{x}') &= k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}') \\k(\mathbf{x}, \mathbf{x}') &= k_3(\phi(\mathbf{x}), \phi(\mathbf{x}')) \\k(\mathbf{x}, \mathbf{x}') &= \mathbf{x}^\top \mathbf{A} \mathbf{x}' \\k(\mathbf{x}, \mathbf{x}') &= k_a(\mathbf{x}_a, \mathbf{x}'_a) + k_b(\mathbf{x}_b, \mathbf{x}'_b) \\k(\mathbf{x}, \mathbf{x}') &= k_a(\mathbf{x}_a, \mathbf{x}'_a)k_b(\mathbf{x}_b, \mathbf{x}'_b)\end{aligned}$$



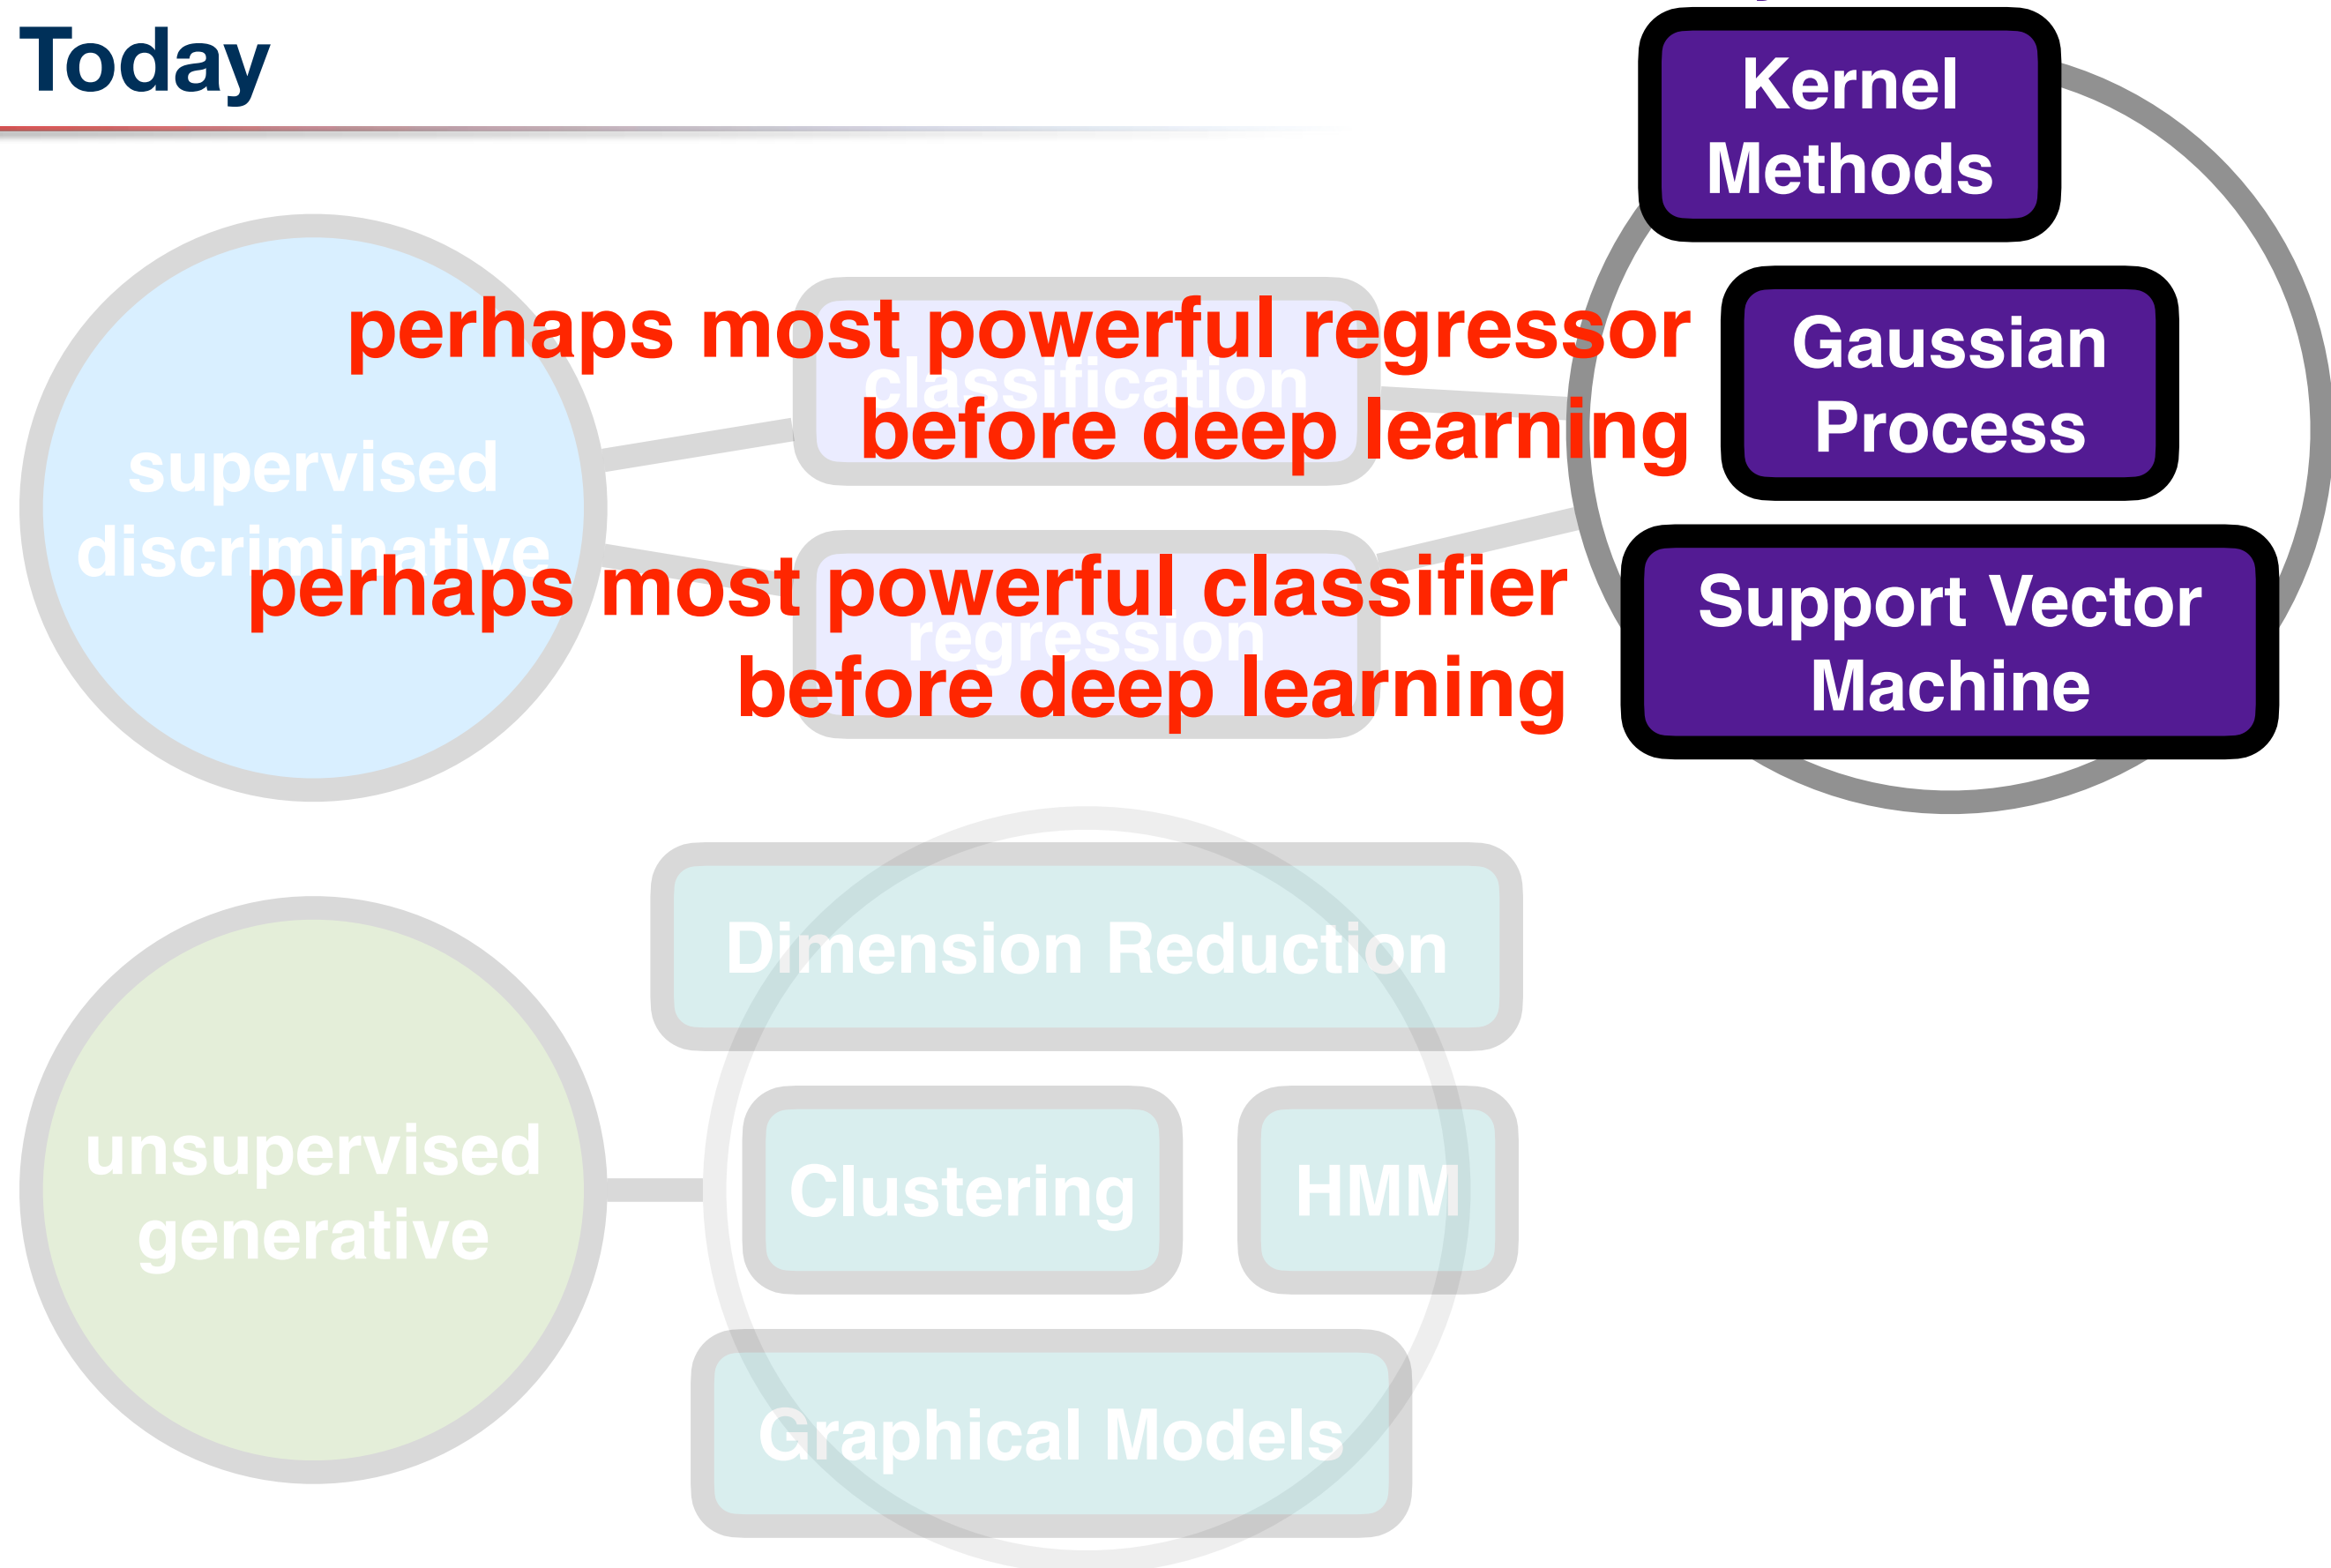
Recap!

- memory-based methods: we want to keep training data points for further use in the prediction
- need to have a metric to evaluate the distance/similarity between data points in the feature space: inner product ➡ kernel!
 - scalar product in x brings possibility to have kernel trick!
- kernel provide a way to compute inner product in feature space
 - simple kernel computation can be related to complicated feature map!
 - the kernel function itself is more important than the feature map (for memory-based methods), we don't even need to know how to compute feature map!
 - usually we imagine kernel is providing us a way to project the data into much higher dimensional space
 - we can easily build kernel from kernel!



Today

memory-based methods



Nonparametric Regression

- Previously on linear parametric regression

$$y(\mathbf{x}, \mathbf{w}) = \mathbf{w}^\top \phi(\mathbf{x})$$

- ▶ prior distribution over \mathbf{w} , e.g., isotropic Gaussian $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$, induces a corresponding distribution over functions $y(\mathbf{x}, \mathbf{w})$
- ▶ while given training samples with specific values $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$, we can get random variables $y(\mathbf{x}_1), y(\mathbf{x}_2), \dots, y(\mathbf{x}_N)$
- ▶ the joint distribution $\mathbf{y} = \{y(\mathbf{x}_1), y(\mathbf{x}_2), \dots, y(\mathbf{x}_N)\}$ is also Gaussian!

$$\mathbb{E}[\mathbf{y}] = \Phi \mathbb{E}[\mathbf{w}] = \mathbf{0}$$

$$\text{cov}[\mathbf{y}] = \mathbb{E}[\mathbf{y}\mathbf{y}^\top] = \Phi \mathbb{E}[\mathbf{w}\mathbf{w}^\top] \Phi^\top = \frac{1}{\alpha} \Phi \Phi^\top = \mathbf{K}$$

Nonparametric Regression

- Previously on linear parametric regression


$$y(\mathbf{x}, \mathbf{w}) = \mathbf{w}^\top \phi(\mathbf{x})$$

- ▶ prior distribution over \mathbf{w} , e.g., isotropic Gaussian $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$, induces a corresponding distribution over functions $y(\mathbf{x}, \mathbf{w})$
- ▶ while given training samples with specific values $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$, we can get random variables $y(\mathbf{x}_1), y(\mathbf{x}_2), \dots, y(\mathbf{x}_N)$
- ▶ the joint distribution $\mathbf{y} = \{y(\mathbf{x}_1), y(\mathbf{x}_2), \dots, y(\mathbf{x}_N)\}$ is also Gaussian!

$$\mathbb{E}[\mathbf{y}] = \Phi \mathbb{E}[\mathbf{w}] = \mathbf{0}$$

$$\text{cov}[\mathbf{y}] = \mathbb{E}[\mathbf{y}\mathbf{y}^\top] = \Phi \mathbb{E}[\mathbf{w}\mathbf{w}^\top] \Phi^\top = \frac{1}{\alpha} \Phi \Phi^\top = \mathbf{K}$$

- Now we are going to introduce “**Gaussian Process**”, which doesn’t use parametric model but instead define a prior probability **distribution over functions** directly.




what do these mean?



Gaussian Process

For any set \mathbf{S} , a Gaussian Process on \mathbf{S} is a set of random variables ($f(x), x \in \mathbf{S}$), so for any $n \in \mathbb{N}$ and $x_1, \dots, x_n \in \mathbf{S}$, $\{f(x_1), \dots, f(x_n)\}$ is (multivariate) Gaussian

- ▶ the number of elements in set \mathbf{S} can be infinite many.
- ▶ the number of random variables $f(x)$ can be any
 - $\{f(x_9), f(x_5), f(x_2), f(x_7)\}$ can build a mean and a covariance function, Gaussian!
 - $\{f(x_7), f(x_8)\}$ can build another mean and another covariance function, Gaussian!



Gaussian Process

For any set \mathbf{S} , a Gaussian Process on \mathbf{S} is a set of random variables $(f(x), x \in \mathbf{S})$, so for any $n \in \mathbb{N}$ and $x_1, \dots, x_n \in \mathbf{S}$, $\{f(x_1), \dots, f(x_n)\}$ is (multivariate) Gaussian

- ▶ the number of elements in set \mathbf{S} can be infinite many.
- ▶ the number of random variables $f(x)$ can be any
 - $\{f(x_9), f(x_5), f(x_2), f(x_7)\}$ can build a mean and a covariance function, Gaussian!
 - $\{f(x_7), f(x_8)\}$ can build another mean and another covariance function, Gaussian!
- ▶ Gaussian Processes (GPs) are parameterised by a mean function $\mu(x)$, and a covariance function, or **kernel**, $K(x, x')$

$$\text{cov}(f_n, f_m) = \langle f_n - \mu(f_n), f_m - \mu(f_m) \rangle = \langle f_n, f_m \rangle - \mu(f_n)\mu(f_m)$$



Gaussian Process

For any set \mathbf{S} , a Gaussian Process on \mathbf{S} is a set of random variables $(f(x), x \in \mathbf{S})$, so for any $n \in \mathbb{N}$ and $x_1, \dots, x_n \in \mathbf{S}$, $\{f(x_1), \dots, f(x_n)\}$ is (multivariate) Gaussian

- ▶ the number of elements in set \mathbf{S} can be infinite many.
- ▶ the number of random variables $f(x)$ can be any
 - $\{f(x_9), f(x_5), f(x_2), f(x_7)\}$ can build a mean and a covariance function, Gaussian!
 - $\{f(x_7), f(x_8)\}$ can build another mean and another covariance function, Gaussian!
- ▶ Gaussian Processes (GPs) are parameterised by a mean function $\mu(x)$, and a covariance function, or **kernel**, $K(x, x')$



[any possible collection of random variables (from training data) creates a function]




Gaussian process is a distribution over functions

ah-ha!



Gaussian Process Regression

- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^N\} = (\mathbf{X}, \mathbf{y})$
 - regression is not a professional statement:
given a bunch of training data, and predict new y for a test input \mathbf{x}
- We want to learn a function f with error bars from data \mathcal{D}




Gaussian process

$$y_n = \boxed{f}(\mathbf{x}_n) + \epsilon_n$$
$$\epsilon_n \sim \mathcal{N}(\cdot | 0, \beta^{-1})$$


Gaussian Process Regression


- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$
 - regression is not professional statement:
given a bunch of training data, and predict new y for a test input \mathbf{x}
- Very basic ideas help you to understand GP regression:
 - It is a memory-based method!!!
 - Prediction is based on the relation between test input and training data!
 - Intuition: if \mathbf{x} and \mathbf{x}' are close to each other (in feature space), then their y will be also close

Gaussian Process Regression



Gaussian Process Regression





Gaussian Process Regression


- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$
 - regression is not a professional statement:
given a bunch of training data, and predict new y for a test input \mathbf{x}
- Very basic ideas help you to understand GP regression:
 - It is a memory-based method!!!
 - Prediction is based on the relation between test input and training data!
 - Intuition: if \mathbf{x} and \mathbf{x}' are close to each other (in feature space), then their y will be also close
 - metric to evaluate the similarity/distance between \mathbf{x} → kernel!
 - GP is parameterised by a mean function $\mu(\mathbf{x})$, and a covariance function, or kernel, $K(\mathbf{x}, \mathbf{x}')$.
 - Since random variable $y_n = f(\mathbf{x}_n)$, we can compute the similarity between random variables by covariance


Gaussian Process Regression

- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$

regression in not professional statement:

The kernel function that determines K should be chosen to express the property that: for points \mathbf{x}_n and \mathbf{x}_m that are similar, the corresponding values $y(\mathbf{x}_n)$ and $y(\mathbf{x}_m)$ will be more strongly correlated than for dissimilar points.

- Intuition: if \mathbf{x} and \mathbf{x}' are close to each other (**in feature space**), then their y will be also close
- metric to evaluate the similarity/distance between \mathbf{x}  kernel!
- GP is parameterised by a mean function $\mu(\mathbf{x})$, and a covariance function, or kernel, $K(\mathbf{x}, \mathbf{x}')$.
- Since random variable $y_n = f(\mathbf{x}_n)$, we can compute the similarity between random variables by covariance



ah-ha!

Gaussian Process Regression

- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$
 - regression is not professional statement:
given a bunch of training data, and predict new y for a test input \mathbf{x}
- Very basic ideas help you to understand GP regression:

$$y_n = f(\mathbf{x}_n) + \epsilon_n$$

$$\epsilon_n \sim \mathcal{N}(\cdot | 0, \beta^{-1})$$

We don't have any parametric model like $y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x})$
Gaussian process regression is nonparametric!

Gaussian Process Regression

- Imagine observing a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$
 - regression is not professional statement:
given a bunch of training data, and predict new y for a test input \mathbf{x}
- Very basic ideas help you to understand GP regression:

$$y_n = f(\mathbf{x}_n) + \epsilon_n$$

$$\epsilon_n \sim \mathcal{N}(\cdot | 0, \beta^{-1})$$


We don't have any parametric model like $y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x})$
Gaussian process regression is nonparametric!



ah-ha!

“nonparametric” doesn’t mean there is no parameters, instead it means: the number of parameters will grow with number of data!

Gaussian Process Regression



$$y_n = f(\mathbf{x}_n) + \epsilon_n$$

$$\epsilon_n \sim \mathcal{N}(\cdot | 0, \beta^{-1})$$

for $\mathbf{f} = [f(\mathbf{x}_1), \dots, f(\mathbf{x}_N)]^\top$ and $\mathbf{y} = [y_1, \dots, y_N]^\top$

$$p(\mathbf{y}|\mathbf{f}) = \mathcal{N}(\mathbf{y}|\mathbf{f}, \beta^{-1}\mathbf{I}_N)$$

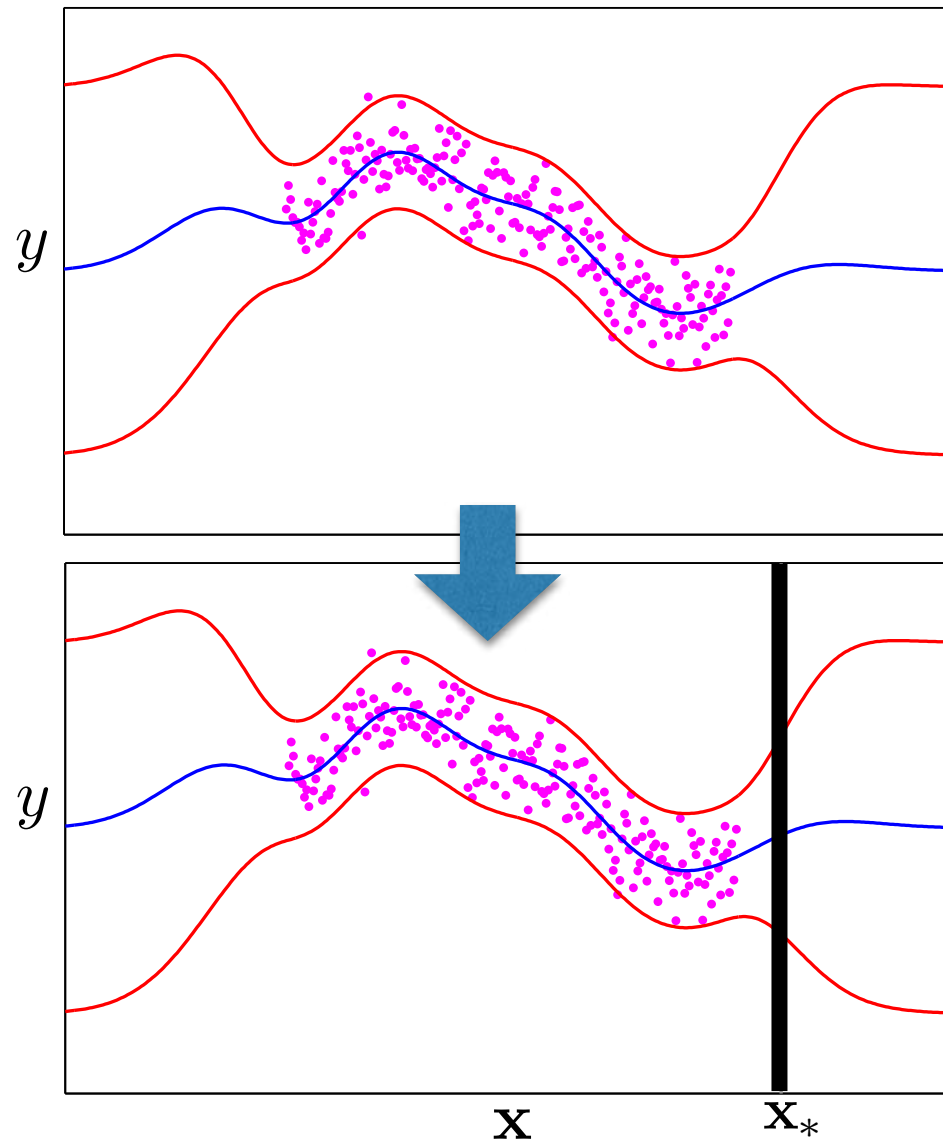
$$p(\mathbf{f}) = \mathcal{N}(\mathbf{0}, \mathbf{K})$$

marginal likelihood
$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C})$$

where the covariance matrix \mathbf{C} has elements

$$\mathbf{C}(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m) + \beta^{-1}\delta_{nm}$$

Gaussian Process Regression



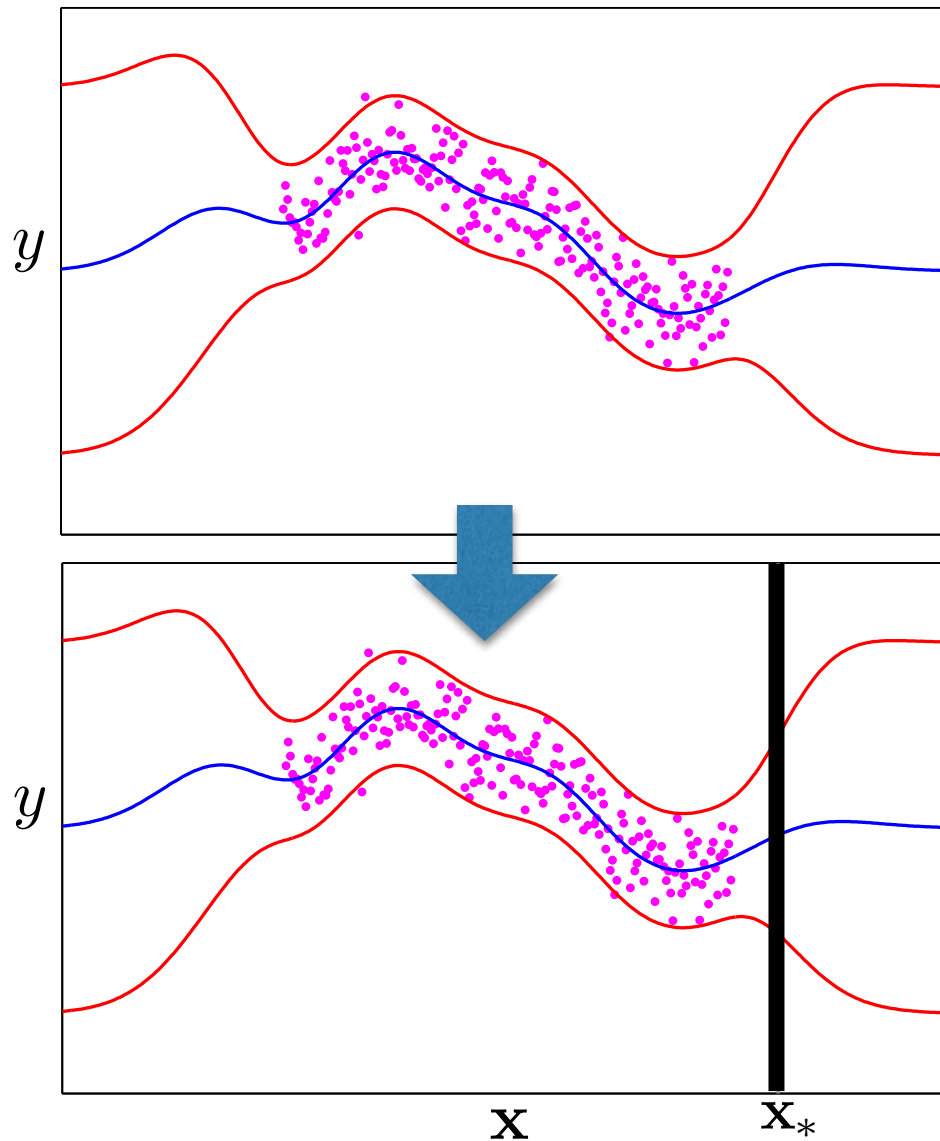
marginal likelihood

$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C})$$

$$\mathbf{C}(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m) + \beta^{-1}\delta_{nm}$$

prediction?

Gaussian Process Regression



marginal likelihood

$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C})$$

$$\mathbf{C}(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m) + \beta^{-1}\delta_{nm}$$


prediction?



ah-ha!

Because we know that all \mathbf{f} (training+testing) together become a multivariate Gaussian distribution \mathbf{G} , therefore if we want to predict the distribution of new \mathbf{f}^* , we just need to compute the covariance matrix of \mathbf{G} , then cut \mathbf{G} on \mathbf{f}^* to see the conditional distribution thus achieve prediction :D

Gaussian Process Regression



marginal likelihood

$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C})$$

$$\mathbf{C}(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m) + \beta^{-1}\delta_{nm}$$

prediction

denote $\mathbf{y}_{N+1} = [\mathbf{y}, y^*]^\top$ and $y^* = f(\mathbf{x}^*)$

$$p(\mathbf{y}_{N+1}) = \mathcal{N}(\mathbf{y}_{N+1}, |\mathbf{0}, \mathbf{C}_{N+1})$$

$$\mathbf{C}_{N+1} = \begin{bmatrix} \mathbf{C} & k(\mathbf{x}, \mathbf{x}^*) \\ k(\mathbf{x}, \mathbf{x}^*)^\top & k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1} \end{bmatrix}$$


conditional distribution $p(y^*|\mathbf{y})$ is a Gaussian distribution with:

$$\mu(\mathbf{x}^*) = k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} \mathbf{y}$$

$$\sigma^2(\mathbf{x}^*) = k^* - k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} k(\mathbf{x}, \mathbf{x}^*)$$

$$k^* = k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1}$$

Gaussian Process Regression



prediction

denote $\mathbf{y}_{N+1} = [\mathbf{y}, y^*]^\top$ and $y^* = f(\mathbf{x}^*)$

$p(\mathbf{y}_{N+1}) = \mathcal{N}(\mathbf{y}_{N+1} | \mathbf{0}, \mathbf{C}_{N+1})$

1° kernel

$$\mathbf{C}_{N+1} = \begin{bmatrix} \mathbf{C} & k(\mathbf{x}, \mathbf{x}^*) \\ k(\mathbf{x}, \mathbf{x}^*)^\top & k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1} \end{bmatrix}$$

conditional distribution $p(y^* | \mathbf{y})$ is a Gaussian distribution with:

2° conditional


$$\mu(\mathbf{x}^*) = k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} \mathbf{y}$$

$$\sigma^2(\mathbf{x}^*) = k^* - k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} k^*$$

3° done!

$$k^* = k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1}$$

Gaussian Process Regression



similarity from testing \mathbf{x}^*
to all the training data!
 $\mathbf{C}^{-1}\mathbf{y}$ is just a vector of scalars

$$\mathbf{C}_{N+1} = \begin{bmatrix} \mathbf{C} & k(\mathbf{x}, \mathbf{x}^*) \\ k(\mathbf{x}, \mathbf{x}^*)^\top & k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1} \end{bmatrix}$$

1^o kernel

conditional distribution $p(y^* | \mathbf{y})$ is a Gaussian distribution with:

2^o conditional

$$\mu(\mathbf{x}^*) = k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} \mathbf{y}$$
$$\sigma^2(\mathbf{x}^*) = k^* - k(\mathbf{x}, \mathbf{x}^*)^\top \mathbf{C}^{-1} k^*$$
$$k^* = k(\mathbf{x}^*, \mathbf{x}^*) + \beta^{-1}$$

3^o done!

Gaussian Process: learning the kernel

Still remember this?

The kernel function that determines K should be chosen to express the property that:
for points \mathbf{x}_n and \mathbf{x}_m that are similar,
the corresponding values $y(\mathbf{x}_n)$ and $y(\mathbf{x}_m)$ will be more strongly correlated than for dissimilar points.

- Consider **covariance function C** with hyper-parameters θ

$$k_{\theta}(\mathbf{x}_n, \mathbf{x}_m) = \theta_0 \exp\left\{-\theta_1 \frac{\|\mathbf{x}_n - \mathbf{x}_m\|^2}{2}\right\} + \theta_2 + \theta_3 \mathbf{x}_n^{\top} \mathbf{x}_m$$

Gaussian Process: learning the kernel

Still remember this?

The kernel function that determines \mathbf{K} should be chosen to express the property that:
for points \mathbf{x}_n and \mathbf{x}_m that are similar,
the corresponding values $y(\mathbf{x}_n)$ and $y(\mathbf{x}_m)$ will be more strongly correlated than for dissimilar points.

- Consider **covariance function** \mathbf{C} with hyper-parameters $\boldsymbol{\theta}$

$$k_{\boldsymbol{\theta}}(\mathbf{x}_n, \mathbf{x}_m) = \theta_0 \exp\left\{-\theta_1 \frac{\|\mathbf{x}_n - \mathbf{x}_m\|^2}{2}\right\} + \theta_2 + \theta_3 \mathbf{x}_n^\top \mathbf{x}_m$$

- Given $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^N\} = (\mathbf{X}, \mathbf{y})$, the marginal likelihood is function of $\boldsymbol{\theta}$

$$p(\mathbf{y}|\boldsymbol{\theta}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C}_{\boldsymbol{\theta}})$$
$$\ln p(\mathbf{y}|\boldsymbol{\theta}) = -\frac{1}{2} \ln |\mathbf{C}_{\boldsymbol{\theta}}| - \frac{1}{2} \mathbf{y}^\top \mathbf{C}_{\boldsymbol{\theta}}^{-1} \mathbf{y} - \frac{N}{2} \ln (2\pi) \quad \text{👉} \quad \frac{\partial \ln p(\mathbf{y}|\boldsymbol{\theta})}{\partial \boldsymbol{\theta}}$$



Gaussian Process: learning the kernel

Still remember this?

The kernel function that determines K should be chosen to express the property that:
for points x_n and x_m that are similar,
the corresponding values $y(x_n)$ and $y(x_m)$ will be more strongly correlated than for dissimilar points.

- Consider **covariance function C** with hyper-parameters θ

The keys in Gaussian Process Regression are:

1. Choose **kernel function**
2. Estimate the proper **hyper-parameters!**

$$p(\mathbf{y}|\theta) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C}_\theta)$$
$$\ln p(\mathbf{y}|\theta) = -\frac{1}{2} \ln |\mathbf{C}_\theta| - \frac{1}{2} \mathbf{y}^\top \mathbf{C}_\theta^{-1} \mathbf{y} - \frac{N}{2} \ln (2\pi) \quad \text{pointing to } \frac{\partial \ln p(\mathbf{y}|\theta)}{\partial \theta}$$



Feature selection can be achieved by ARD in GPs

Problem: Often there are *many* possible inputs that might be relevant to predicting a particular output. We need algorithms that automatically decide which inputs are relevant.

Automatic Relevance Determination:

Consider this covariance function:

$$\mathbf{K}_{nn'} = v \exp \left[-\frac{1}{2} \sum_{d=1}^D \left(\frac{x_n^{(d)} - x_{n'}^{(d)}}{r_d} \right)^2 \right]$$

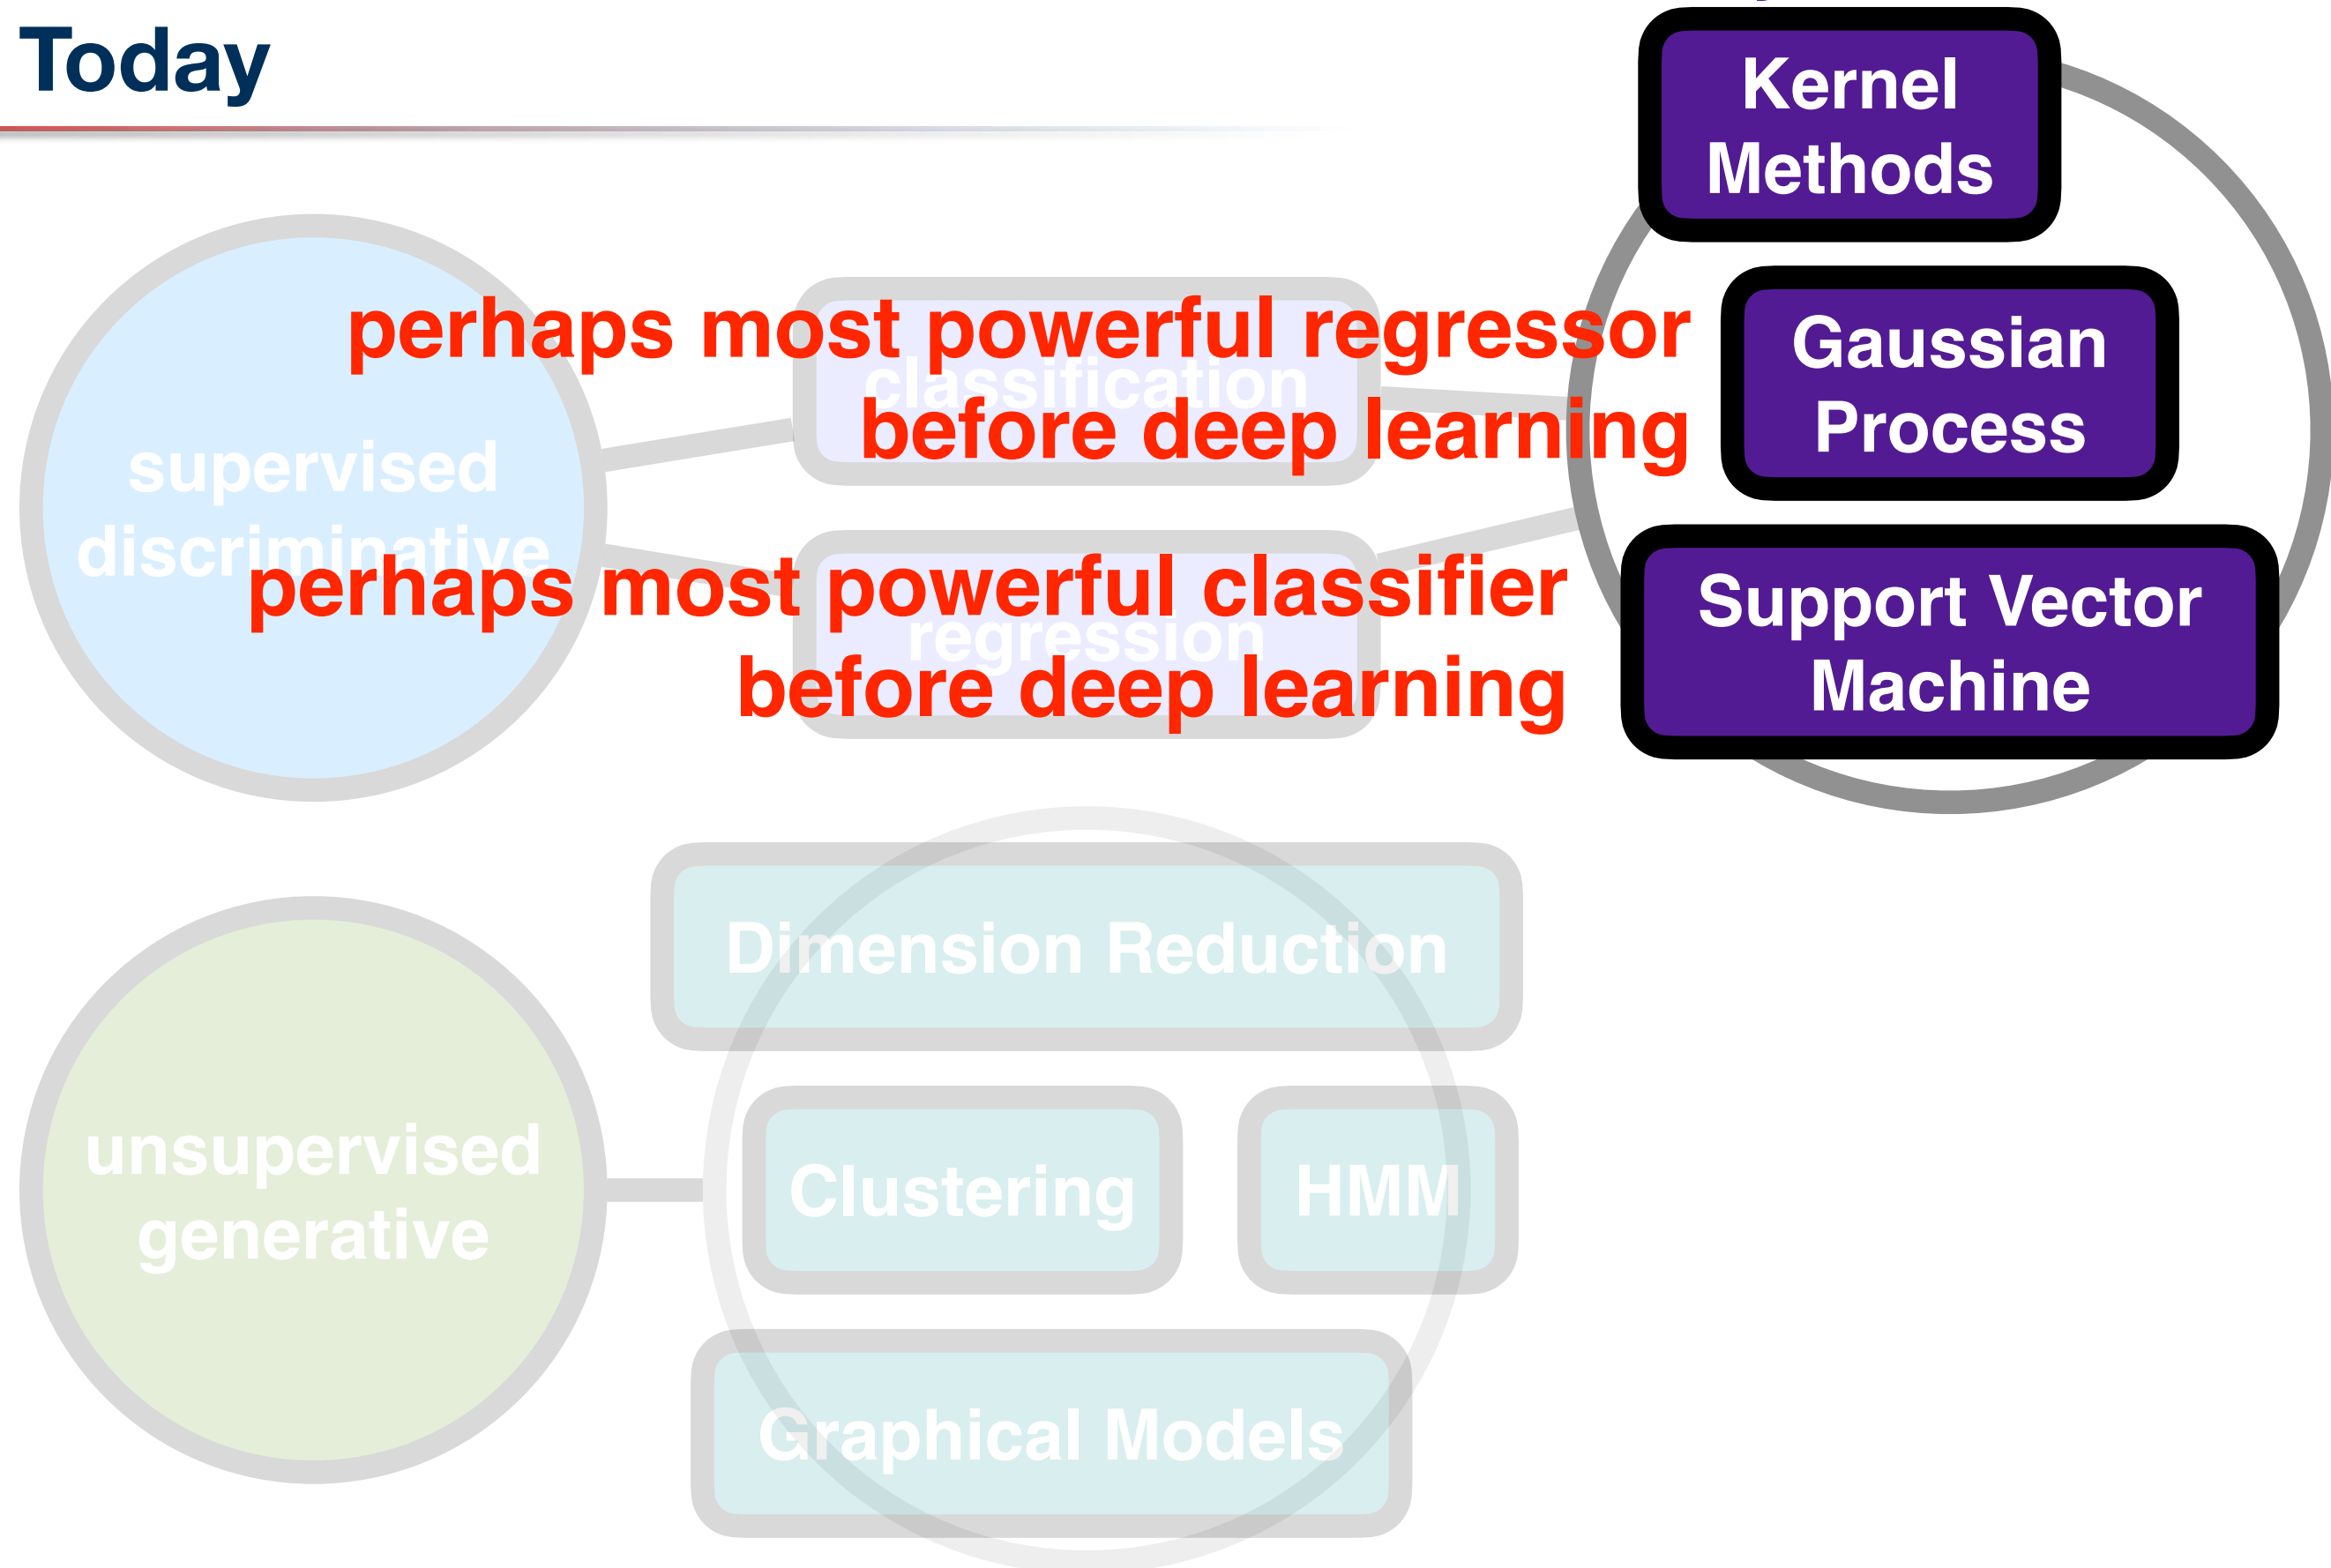
The parameter r_d is the **length scale of the function along input dimension d** .

As $r_d \rightarrow \infty$ the function f varies less and less as a function of $x^{(d)}$, that is, the d th dimension becomes *irrelevant*.

Given data, by learning the lengthscales (r_1, \dots, r_D) it is possible to do automatic feature selection.

Today

memory-based methods



Sparse kernel machines


- Memory based methods: we would like to utilise the training data in the prediction!
 - What we have introduced (KNN, Parzen probability estimation, GPs) need to keep “ALL” training data ... heavy memory demands!
 - Support Vector Machine (SVM), what we are going to introduce, is a kernel-based algorithm that have **sparse solutions**!
 - Predictions by kernel functions evaluated on **subset** of training data!

Sparse kernel machines

- Let's start from simple **binary classification** problem:

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

- classification based on sign of $y(\mathbf{x})$
- please note that explicit computing on $\phi(\mathbf{x})$ will be avoided by dual form
👉 kernel!

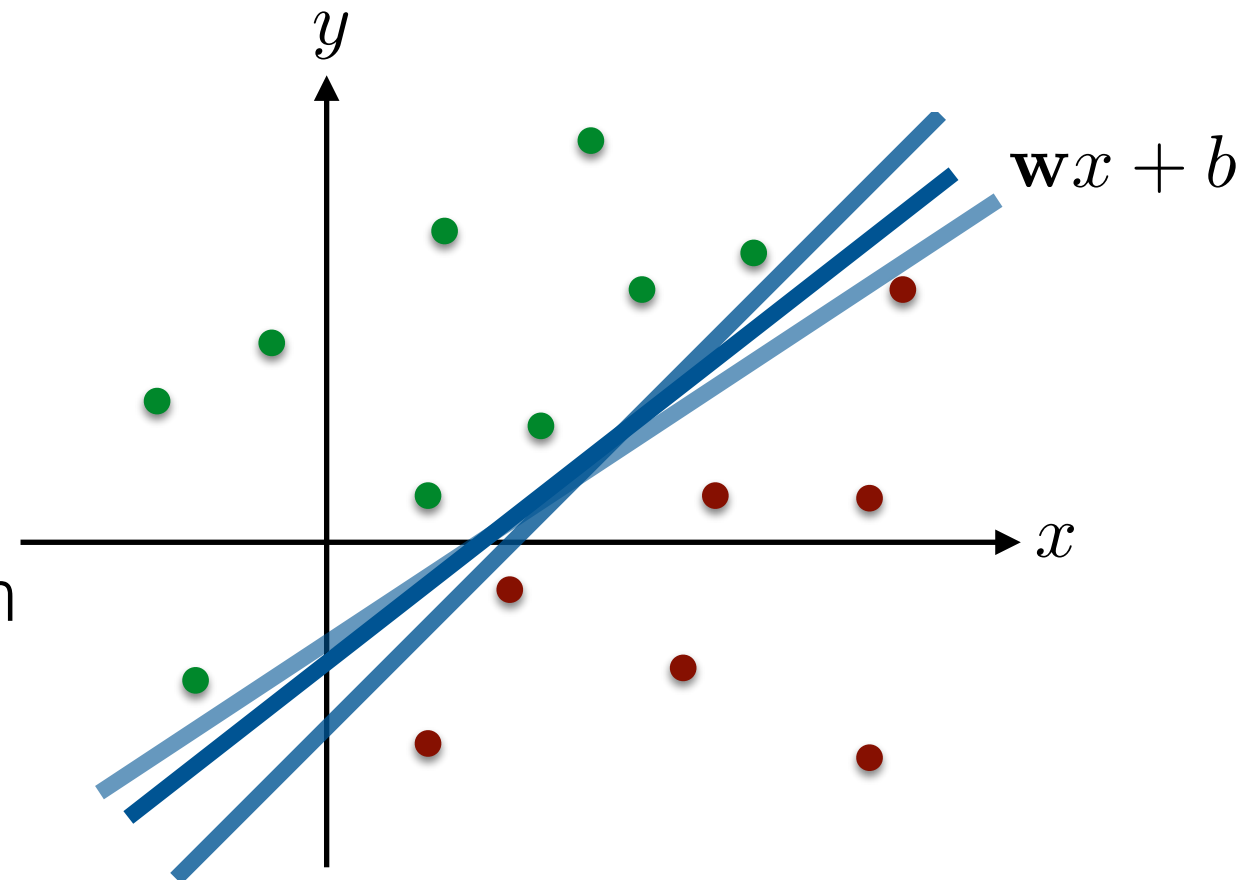


Sparse kernel machines


- Let's start from simple **binary classification** problem:

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

- classification based on sign of $y(\mathbf{x})$
- please note that explicit computing on $\phi(\mathbf{x})$ will be avoided by dual form
👉 kernel!



- When assuming the data is linearly separable, it means there will be infinite solutions for \mathbf{w} and b , we should try to find the one with minimum generalisation error.




what does it mean?

Sparse kernel machines

- Let's start from simple **binary classification** problem:

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$


Maximise margin!
margin is the smallest distance b/w decision boundary and any of the training samples



- When assuming the data is linearly separable, it means there will be infinite solutions for \mathbf{w} and b , we should try to find the one with minimum generalisation error.

Sparse kernel machines

Maximise margin!
margin is the smallest
distance b/w decision
and any of the training



$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

perpendicular distance from a data point x to a hyperplane $y(\mathbf{x}) = 0$ is given by

$$\frac{|y(\mathbf{x})|}{\|\mathbf{w}\|}$$

Sparse kernel machines

we are only interested in solutions for which all data points are correctly classified
so that $t_n y(\mathbf{x}_n) > 0 \quad \forall n$

t_n and $y(\mathbf{x}_n)$ with
the same sign

$$\frac{t_n y(\mathbf{x})}{\|\mathbf{x}\|} = \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{x}\|}$$

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

perpendicular distance from a data point x to a hyperplane $y(\mathbf{x}) = 0$ is given by

$$\frac{|y(\mathbf{x})|}{\|\mathbf{w}\|} \text{ absolute value}$$

Sparse kernel machines

we are only interested in solutions for which all data points are correctly classified
so that $t_n y(\mathbf{x}_n) > 0 \quad \forall n$

$$\frac{t_n y(\mathbf{x})}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}$$

we want to maximise the margin!

$$\underset{\mathbf{w}, b}{\operatorname{argmax}} \left\{ \min_n \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|} \right\}$$

**margin is the smallest
distance b/w decision boundary
and any of the training samples**

Sparse kernel machines

we are only interested in solutions for which all data points are correctly classified
so that $t_n y(\mathbf{x}_n) > 0 \quad \forall n$

$$\frac{t_n y(\mathbf{x})}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}$$

we want to maximise the margin!

$$\operatorname{argmax}_{\mathbf{w}, b} \left\{ \min_n \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|} \right\}$$

note here if we scale \mathbf{w}, b together $\frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}$ won't change

so we set a condition to get the unique solution:

for the points closet to the hyperplane, $t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) = 1$

which means, for all data points, $t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1$



Sparse kernel machines

- Add these two criteria together!

$$\operatorname{argmax}_{\mathbf{w}, b} \left\{ \min_n \frac{t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|} \right\} \quad t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) = 1$$

- The overall optimisation problem becomes:

$$\begin{aligned} & \text{maximise } \frac{1}{\|\mathbf{w}\|} \Rightarrow \text{minimise } \|\mathbf{w}\| \\ & \text{subject to } t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1 \end{aligned}$$

- Here comes Lagrange multipliers $a_n \geq 0$ $\mathbf{a} = (a_1, \dots, a_N)^\top$

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_{n=1}^N a_n \{t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1\}$$

remember regularisation story?

- minimise w.r.t. \mathbf{w} and b while maximise w.r.t. \mathbf{a}

Sparse kernel machines

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_{n=1}^N a_n \{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1\}$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial \mathbf{w}} = 0 \quad \text{👉} \quad \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial b} = 0 \quad \text{👉} \quad 0 = \sum_{n=1}^N a_n t_n$$



ah-ha!

w is the linear combination based on the feature representations of training data.
If you put this w back to L(w, b, a), it is easier to see that L has all x appearing as scalar product! 👉 kernel trick!

Sparse kernel machines

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_{n=1}^N a_n \{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1\}$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial \mathbf{w}} = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

cancel out

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial b} = 0 \quad \Rightarrow \quad 0 = \sum_{n=1}^N a_n t_n$$

- put \mathbf{w} back to $L(\mathbf{w}, b, \mathbf{a})$ to get the **dual representation of the max-margin problem**, in which we maximise:

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

subject to $a_n \geq 0, \quad n = 1, \dots, N$

$$\sum_{n=1}^N a_n t_n = 0$$

Sparse kernel machines

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

subject to $a_n \geq 0, \quad n = 1, \dots, N$

$$\sum_{n=1}^N a_n t_n = 0$$

use SMO (Sequential Minimal optimisation)
for QP (Quadratic Programming) to solution of a
objective is quadratic and so any local optimum
will also be a global optimum provided that the
constraints define a convex region (since being linear)

Sparse kernel machines

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

$$\text{subject to } a_n \geq 0, \quad n = 1, \dots, N$$

$$\sum_{n=1}^N a_n t_n = 0$$

after getting the solution for \mathbf{a} , the prediction can be done by:

$$\begin{aligned} y(\mathbf{x}) &= \mathbf{w}^\top \phi(\mathbf{x}) + b \\ \mathbf{w} &= \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \end{aligned}$$

$$\Rightarrow y(\mathbf{x}) = \sum_{n=1}^N a_n t_n \underbrace{k(\mathbf{x}, \mathbf{x}_n)}_{\text{kernel!}} + b$$

memory-based method!

Support Vector Machine

- See from **KKT condition**

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_{n=1}^N a_n \{ t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1 \}$$

$$a_n \geq 0$$

$$t_n y(\mathbf{x}_n) - 1 \geq 0$$

$$a_n \{ t_n y(\mathbf{x}_n) - 1 \} = 0$$

for every data point, either $a_n = 0$ or $t_n y(\mathbf{x}_n) = 1$


data points having this condition
won't be used for prediction

data points having this condition
are called: **support vector**

prediction: $y(\mathbf{x}) = \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) + b$

Support Vector Machine

These data points (vectors) support the hyperplane! we can call them support vectors!



$$\underline{t_n y(\mathbf{x}_n) = 1}$$

data points having this condition are called: **support vector**

remember few slides ago?


note here if we scale \mathbf{w}, b together $\frac{t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}$ won't change

so we set a condition to get the unique solution:

for the points closet to the hyperplane, $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) = 1$


which means, for all data points, $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1$

Support Vector Machine




linearly separable in feature space
(whose dimensionality could be
much higher due to kernel function)

how decision boundary could
be in original data space



KKT condition

- simple optimisation problem
 - ▶ Green line represents equality constraint $\mathbf{h}(\mathbf{x},\mathbf{y})=\mathbf{c}$
 - Blue line shows contour lines of function \mathbf{f}
 - ▶ Intuitively you can see in the optimum, the gradient of \mathbf{f} is parallel to the gradient of \mathbf{h}



$$\text{let } L(X,\lambda,\mu)=f(X)+\sum_{j=1}^p \lambda_j \underbrace{h_j(X)}_{\text{equality constraints}}+\sum_{k=1}^q \mu_k \underbrace{g_k(X)}_{\text{inequality constraints}}$$

$$\nabla f \in \text{span}\{\nabla h\}$$

$$\text{i.e. } \nabla f = -\eta \nabla h$$

equality constraints inequality constraints

KKT condition

$$\left. \frac{\partial L}{\partial X} \right|_{X=X^*} = 0$$

$$\lambda_j \neq 0,$$

$$u_k \geq 0,$$

$$u_k g_k(X^*) = 0$$

$$h_j(X^*) = 0 \quad j=1,2,\dots,p$$

$$g_k(X^*) \leq 0 \quad k=1,2,\dots,q$$

equality constraints should be satisfied

This holds for both interior or boundary solution, called as complementary slackness. (for interior solution, inequality constraints inactive, $u=0$)

inequality constraints to be active

Support Vector Machine - with Soft Margin!

- Data points are allowed to be on the “wrong side” of boundary
 - But with a penalty ~~inversely~~ proportional to the distance from boundary

slack: $\xi_n = |t_n - y(\mathbf{x}_n)|$


- inequality $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1$ replaced by $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1 - \xi_n$

also $\xi_n \geq 0$

- optimisation problem now as
subject to two constraints above

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

also minimise slack



Support Vector Machine - with Soft Margin!

- Data points are allowed to be on the “wrong side” of boundary
 - But with a penalty inversely proportional to the distance from boundary

slack: $\xi_n = |t_n - y(\mathbf{x}_n)|$

- inequality $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1$ replaced by $t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1 - \xi_n$

also $\xi_n \geq 0$

- optimisation problem now as
subject to two constraints above

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

also minimise slack

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

- KKT condition

$$a_n \geq 0$$

$$t_n y(\mathbf{x}_n) - 1 + \xi_n \geq 0$$

$$a_n (t_n y(\mathbf{x}_n) - 1 + \xi_n) = 0$$

$$\mu_n \geq 0$$

$$\xi_n \geq 0$$

$$\mu_n \xi_n = 0$$

Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

- KKT condition

a_n	≥ 0	$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow$	$\mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$
$t_n y(\mathbf{x}_n) - 1 + \xi_n$	≥ 0		
$a_n (t_n y(\mathbf{x}_n) - 1 + \xi_n)$	$= 0$	$\frac{\partial L}{\partial b} = 0 \Rightarrow$	$\sum_{n=1}^N a_n t_n = 0$
μ_n	≥ 0		
ξ_n	≥ 0	$\frac{\partial L}{\partial \xi_n} = 0 \Rightarrow$	$a_n = C - \mu_n$
$\mu_n \xi_n$	$= 0$		

Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

- KKT condition

$$\begin{array}{ll}
 \boxed{a_n \geq 0} & \frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \\
 t_n y(\mathbf{x}_n) - 1 + \xi_n \geq 0 & \\
 a_n (t_n y(\mathbf{x}_n) - 1 + \xi_n) = 0 & \frac{\partial L}{\partial b} = 0 \Rightarrow \sum_{n=1}^N a_n t_n = 0 \\
 \boxed{\mu_n \geq 0} & \frac{\partial L}{\partial \xi_n} = 0 \Rightarrow \boxed{a_n = C - \mu_n} \\
 \xi_n \geq 0 & \\
 \mu_n \xi_n = 0 & \begin{array}{l} \text{👉 } a_n \leq C \\ \text{👉 } 0 \leq a_n \leq C \end{array}
 \end{array}$$

Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

put back to $L(\mathbf{w}, b, \mathbf{a})$

$$\text{⤵} \quad \tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

subject to $0 \leq a_n \leq C$

$$\sum_{n=1}^N a_n t_n = 0$$

use SMO for QP to get solution of a



Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

likewise, after getting the solution for \mathbf{a} , the prediction can be done by:

$$\begin{aligned} y(\mathbf{x}) &= \mathbf{w}^\top \phi(\mathbf{x}) + b \\ \mathbf{w} &= \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \end{aligned}$$

☞ $y(\mathbf{x}) = \sum_{n=1}^N a_n t_n \underline{k(\mathbf{x}, \mathbf{x}_n)} + b$
kernel!

memory-based method!



Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

- KKT condition

$$\begin{array}{rcl} a_n & \geq & 0 \\ t_n y(\mathbf{x}_n) - 1 + \xi_n & \geq & 0 \\ a_n (t_n y(\mathbf{x}_n) - 1 + \xi_n) & = & 0 \\ \mu_n & \geq & 0 \\ \xi_n & \geq & 0 \\ \mu_n \xi_n & = & 0 \end{array}$$

for every data point,
either $a_n = 0$ or $t_n y(\mathbf{x}_n) = 1 - \xi_n$

**not used in
prediction**

**data points related
to support vector**

prediction: $y(\mathbf{x}) = \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) + b$




Support Vector Machine - with Soft Margin!

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

where $a_n \geq 0, \mu_n \geq 0$ are Lagrange multipliers

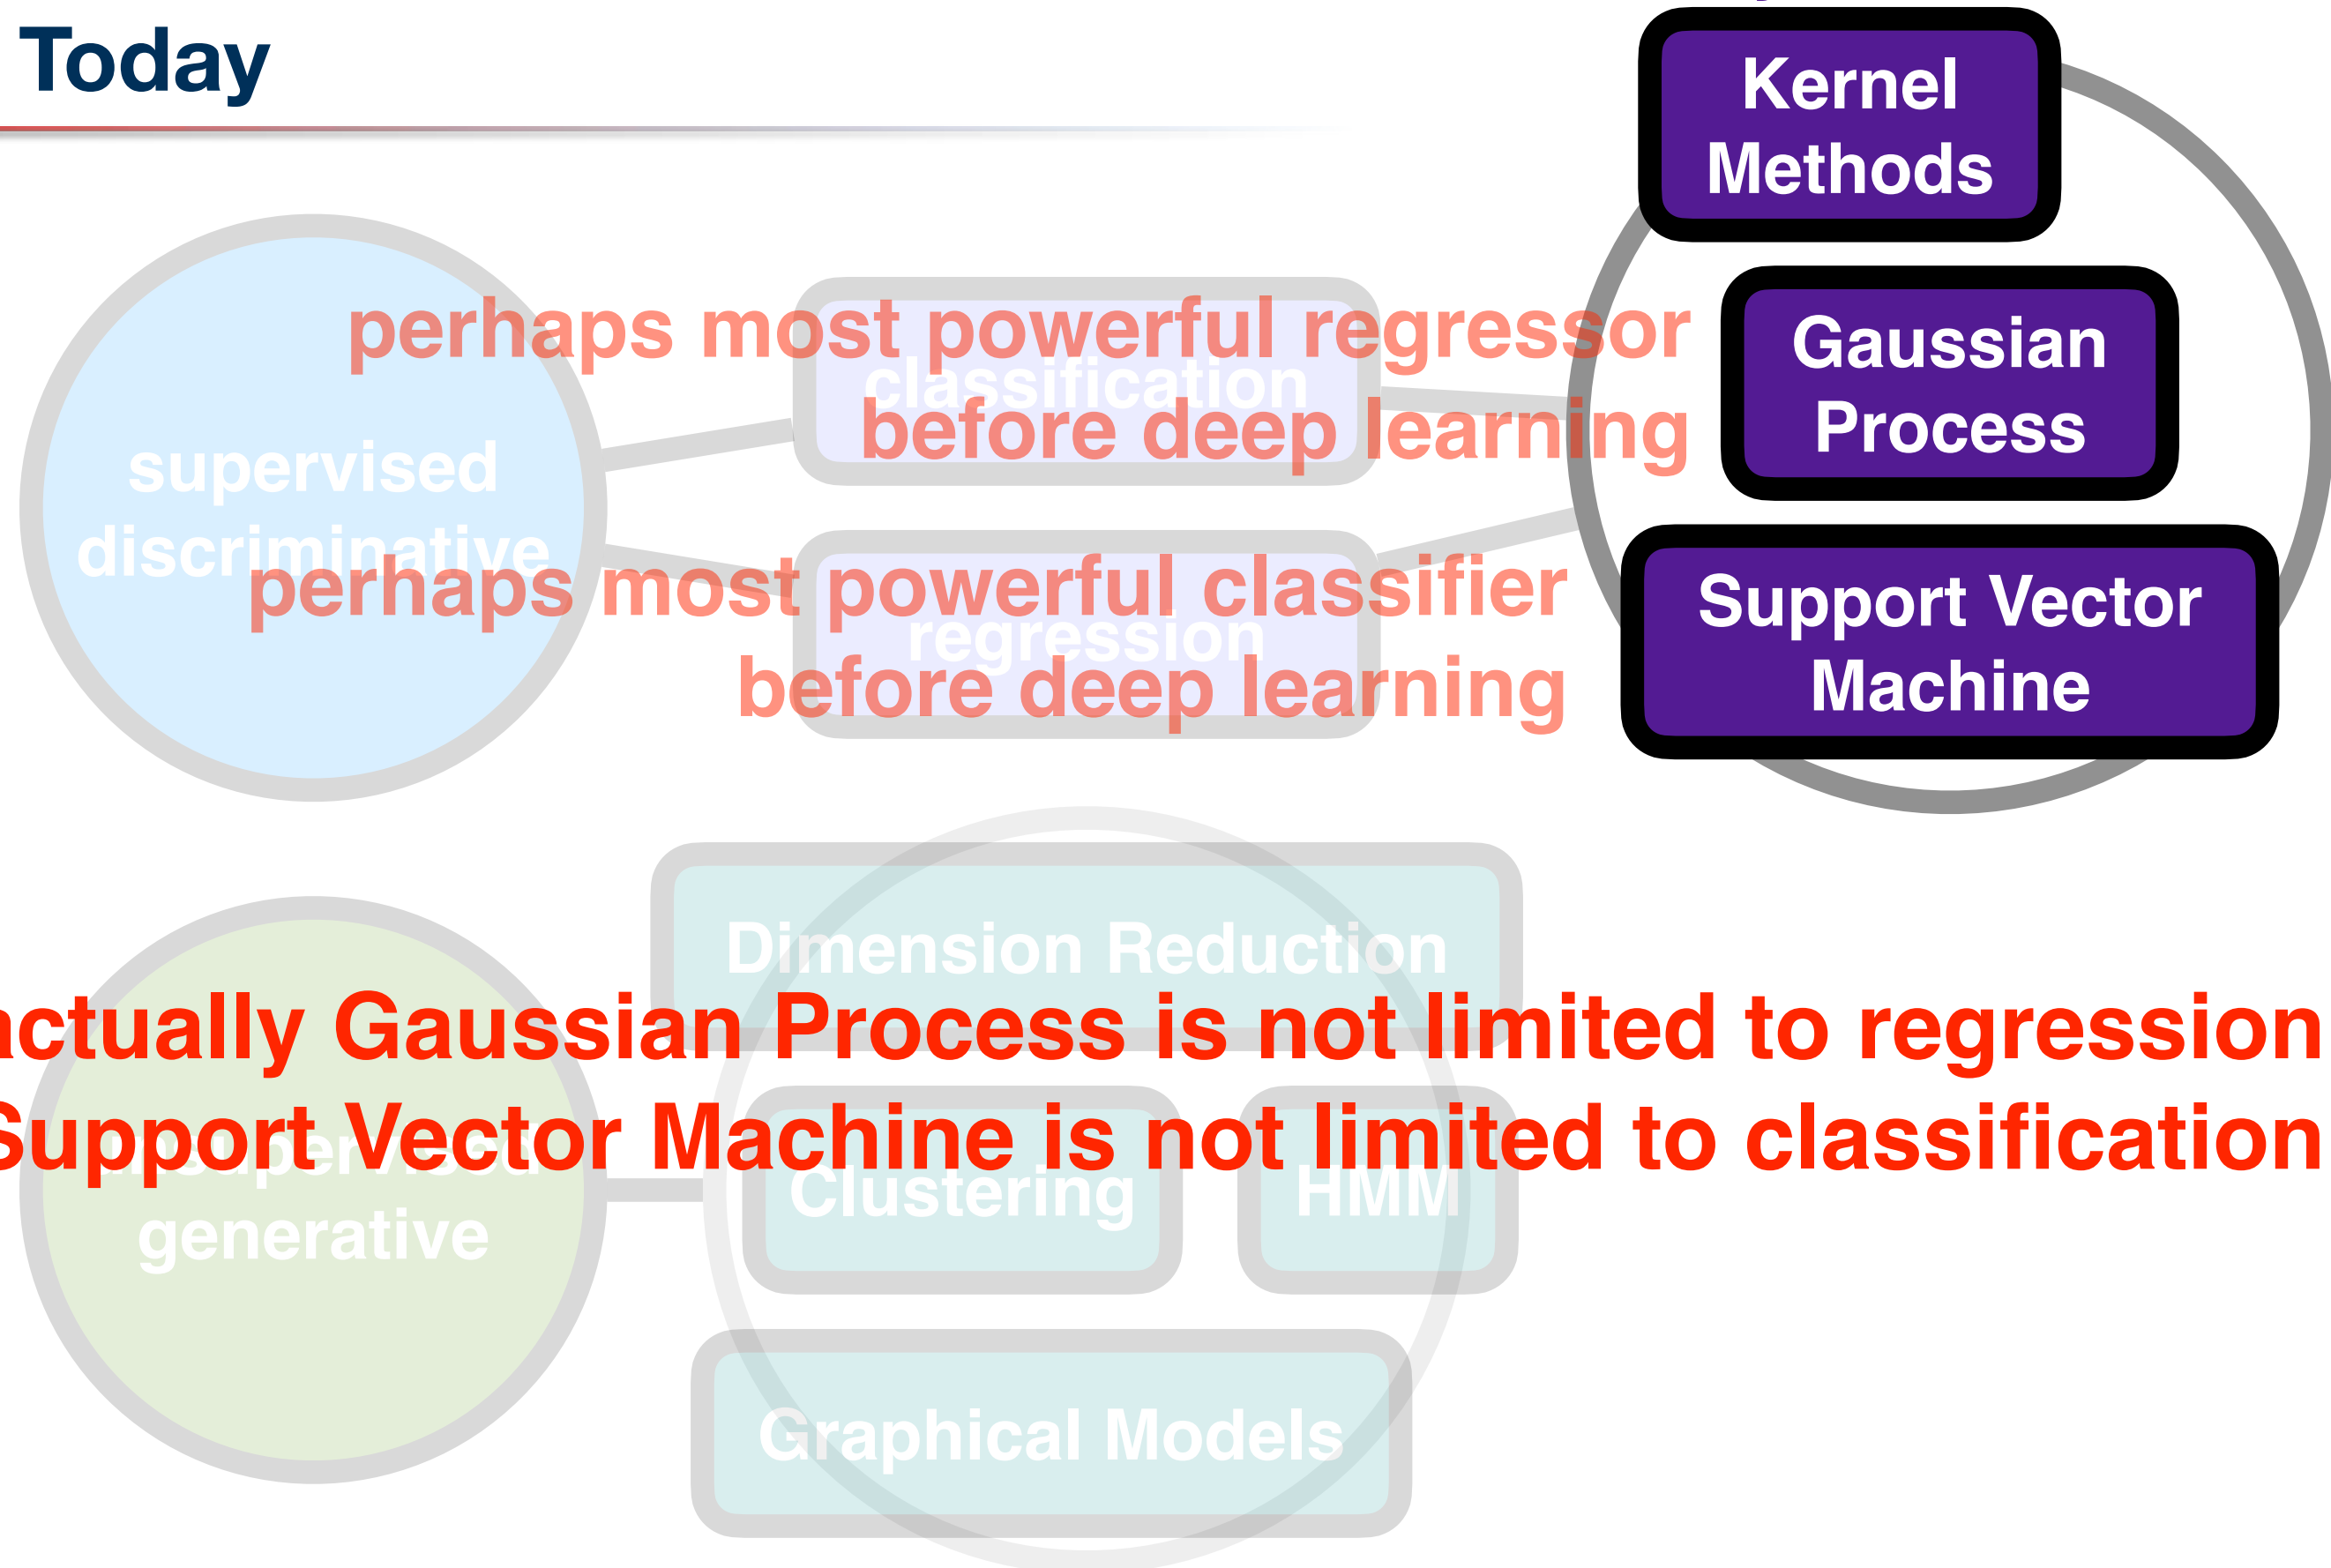
$$0 \leq a_n \leq C$$

- if $a_n = 0$, then such points not used in prediction
- if $a_n < C$ because $a_n = C - \mu_n$ implies $\mu_n > 0$, from $\mu_n \xi_n = 0$ we get $\xi_n = 0$
- if $a_n = C$, then such points lie inside the margin



Today

memory-based methods



Support Vector Machine for Regression


- In simple linear regression, we minimise a regularised error function given by

$$\frac{1}{2} \sum_{n=1}^N \{y_n - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- To obtain sparse solution, we define an **ϵ -insensitive error function** which gives zero error if the absolute difference between the prediction $y(\mathbf{x})$ and the target t is less than ϵ where $\epsilon > 0$

- Example:

$$E_{\epsilon}(y(\mathbf{x}) - t) = \begin{cases} 0, & \text{if } |y(\mathbf{x}) - t| < \epsilon; \\ |y(\mathbf{x}) - t| - \epsilon, & \text{otherwise} \end{cases}$$




→ $C \sum_{n=1}^N E_{\epsilon}(y(\mathbf{x}_n) - t_n) + \frac{1}{2} \|\mathbf{w}\|^2$

Support Vector Machine for Regression

- With also slack :)

Illustration of SVM regression, showing the regression curve together with the ϵ -insensitive 'tube'. Also shown are examples of the slack variables ξ and $\hat{\xi}$. Points above the ϵ -tube have $\xi > 0$ and $\hat{\xi} = 0$, points below the ϵ -tube have $\xi = 0$ and $\hat{\xi} > 0$, and points inside the ϵ -tube have $\xi = \hat{\xi} = 0$.



$$C \sum_{n=1}^N (\xi_n + \hat{\xi}_n) + \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to } \xi_n \geq 0 \quad \hat{\xi}_n \geq 0$$


$$t_n \leq y(\mathbf{x}_n) + \epsilon + \xi_n$$

$$t_n \geq y(\mathbf{x}_n) - \epsilon - \hat{\xi}_n$$

Support Vector Machine for Regression

- With also slack :)

Illustration of SVM regression, showing the regression curve together with the ϵ -insensitive 'tube'. Also shown are examples of the slack variables ξ and $\hat{\xi}$. Points above the ϵ -tube have $\xi > 0$ and $\hat{\xi} = 0$, points below the ϵ -tube have $\xi = 0$ and $\hat{\xi} > 0$, and points inside the ϵ -tube have $\xi = \hat{\xi} = 0$.



$$\begin{aligned}
 L = & C \sum_{n=1}^N (\xi_n + \hat{\xi}_n) + \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N (\mu_n \xi_n + \hat{\mu}_n \hat{\xi}_n) \\
 & - \sum_{n=1}^N a_n (\epsilon + \xi_n + y_n - t_n) - \sum_{n=1}^N \hat{a}_n (\epsilon + \hat{\xi}_n - y_n + t_n)
 \end{aligned}$$

Support Vector Machine for Regression

- Dual representation
 - Maximising

$$\begin{aligned}\tilde{L}(\mathbf{a}, \hat{\mathbf{a}}) = & -\frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N (a_n - \hat{a}_n)(a_m - \hat{a}_m) k(\mathbf{x}_n, \mathbf{x}_m) \\ & -\epsilon \sum_{n=1}^N (a_n + \hat{a}_n) + \sum_{n=1}^N (a_n - \hat{a}_n) t_n\end{aligned}$$

$$\begin{aligned}\text{subject to } & 0 \leq a_n \leq C \\ & 0 \leq \hat{a}_n \leq C\end{aligned}$$

$$\sum_{n=1}^N (a_n - \hat{a}_n) = 0$$

prediction: $y(\mathbf{x}) = \sum_{n=1}^N (a_n - \hat{a}_n) k(\mathbf{x}, \mathbf{x}_n) + b$

Support Vector Machine for Regression


- KKT condition

$$a_n(\epsilon + \xi_n + y_n - t_n) = 0$$

$$\hat{a}_n(\epsilon + \hat{\xi}_n - y_n + t_n) = 0$$

$$(C - a_n)\xi_n = 0$$



$$(C - \hat{a}_n)\hat{\xi}_n = 0.$$



1. a_n can only be nonzero if $\epsilon + \xi_n + y_n - t_n = 0$
 ➡ data points either lies on the upper boundary of the tube: $\xi_n = 0$, or above $\xi_n > 0$
2. \hat{a}_n can only be nonzero if $\epsilon + \hat{\xi}_n - y_n + t_n = 0$
 ➡ data points either lies on the lower boundary of the tube: $\hat{\xi}_n = 0$, or below $\hat{\xi}_n > 0$
3. $\epsilon + \xi_n + y_n - t_n$ and $\epsilon + \hat{\xi}_n - y_n + t_n$ are incompatible
 ➡ for every data point, either a_n or \hat{a}_n (or both) must be nonzero
4. Support vectors are those data points either a_n nonzero or \hat{a}_n nonzero
 ➡ these are data points that lie on the boundary of the tube or outside the tube
5. All points within the tube have $a_n = \hat{a}_n = 0$

Gaussian Process for Classification

- Binary classification problem
 - Given dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i)_{i=1}^n\} = (X, \mathbf{y})$ where \mathbf{y} are binary labels $\{-1, +1\}$, infer class label probabilities at new points.



- There are many ways to relate function values to class probabilities

$$p(y_i|f_i) = \begin{cases} \frac{1}{1+\exp(-y_i f_i)} \\ \Phi(y_i f_i) \\ \mathbf{H}(y_i f_i) \\ \epsilon + (1 - 2\epsilon)\mathbf{H}(y_i f_i) \end{cases}$$

sigmoid (logistic)

cumulative normal (probit)

threshold

robust threshold

- Inference not that easy: approximation.

Limitations of SVM w.r.t. GPs

- The outputs of an SVM represent decisions rather than posterior probabilities.
- The SVM was originally formulated for two classes, and the extension to $K > 2$ classes is problematic.
- Gaussian Process can learn the kernel parameters automatically from data, no cross-validation is needed.
- Gaussian Process can be used for automatic feature selection.
- Gaussian Process can incorporate interpretable noise models and priors over functions, and can sample from prior to get intuitions about the model assumptions.