

Machine Learning Homework 7

K-means & spectral clustering



數據所碩一葉詠富 310554031

Github: <https://github.com/frankye1000/NYCU-MachineLearning/tree/master/HW7>





I. Kernel Eigenfaces

- Part1: Use PCA and LDA to show the first 25 eigenfaces and fisherfaces, and randomly pick 10 images to show their reconstruction. (please refer to the lecture slides).

- First 25 eigenfaces & fisherfaces

| eigenfaces | fisherfaces |
|--|---|
|  |  |

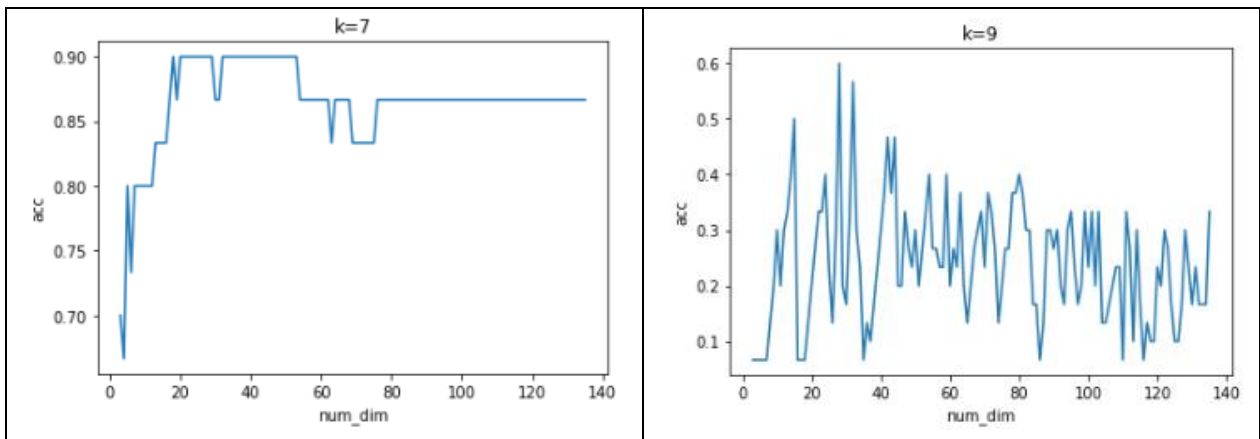
- Random pick 10 images recostruction

| eigenfaces | fisherfaces |
|---|---|
| <p>Original:</p>  <p>Reconstruction:</p>  | <p>Original:</p>  <p>Reconstruction:</p>  |

- Part2: Use PCA and LDA to do face recognition, and compute the performance. You should use k nearest neighbor to classify which subject the testing image belongs to.

- I try num_dim:3~135, k: 3~9 to get the best performance.

| PCA | LDA |
|--------------------------|--------------------------|
| num_dim:18, k:7, acc:90% | num_dim:28, k:9, acc:60% |



➤ Part3: Use kernel PCA and kernel LDA to do face recognition, and compute the performance. (You can choose whatever kernel you want, but you should try different kernels in your implementation.) Then compare the difference between simple LDA/PCA and kernel LDA/PCA, and the difference between different kernels.

- In order to compare the difference between the simple LDA and the kernel LDA, I did not change the parameters num_dim & k.
- We can find the kernel PCA have a little worse performance than simple PCA. The RBF kernel is better than polynomial kernel.
- We can find the kernel LDA have the worse performance than simple LDA. The RBF kernel is better than polynomial kernel.

| kernel | kernel PCA | kernel LDA |
|------------|-----------------------------|-----------------------------|
| RBF | num_dim:18, k:7, acc:86.67% | num_dim:28, k:9, acc:56.67% |
| polynomial | num_dim:18, k:7, acc:83.33% | num_dim:28, k:9, acc:46.67% |

Code:

1. Read image.

```
def read_image(path, H, W):
    pics = os.listdir(path)
    images = np.zeros((W*H, len(pics)))
    labels = np.zeros(len(pics)).astype('uint8')
    for pic, i in zip(pics, np.arange(len(pics))):
        labels[i] = int(pic.split('.')[0][7:9])-1
        image = np.asarray(Image.open(os.path.join(path, pic)).resize((W,H), Image.ANTIALIAS)).flatten()
        images[:,i] = image

    return images, labels
```

2. Print face & reconstruction result.

```

def print_face(X, num, H, W):
    print("fisherface: ")
    n = int(num**0.5)
    axes = []
    fig = plt.figure()
    for i in range(num):
        axes.append(fig.add_subplot(n, n, i+1) )
        axes[-1].axis('off')
        plt.imshow(X[:,i].reshape(H,W), cmap='gray')

    fig.tight_layout()
    plt.show()

def print_reconstruction(X_original, X_reconstruction, num, H, W):
    print("Original: ")
    randint = np.random.choice(X_original.shape[1], num)
    axes = []
    fig = plt.figure()
    for i in range(num):
        axes.append(fig.add_subplot(1, num, i+1) )
        axes[-1].axis('off')
        plt.imshow(X_original[:,randint[i]].reshape(H,W), cmap='gray')
    fig.tight_layout()
    plt.show()

    print("Reconstruction: ")
    axes = []
    fig = plt.figure()
    for i in range(num):
        axes.append(fig.add_subplot(1, num, i+1) )
        axes[-1].axis('off')
        plt.imshow(X_reconstruction[:,randint[i]].reshape(H,W), cmap='gray')

    fig.tight_layout()
    plt.show()

```

3. PCA

I use the concept of SVD($U @ \text{Sigma} @ V.T$) to calculate PCA.

Theorem:
 $A: m \times n$, 则 A 可作 SVD

Note:
 $A = U \Sigma V^T: m \times n$ SVD
 (1) $A^T A = (U \Sigma V^T)^T (U \Sigma V^T)$ (秩 Σ, V)
 $A^T A = V \Sigma^T U^T U \Sigma V^T$
 $= V \Sigma^T \Sigma V^T$ (PDP)
 $A^T A V = V \Sigma^T \Sigma$
 $\Rightarrow A^T A [\vec{v}_1 \dots \vec{v}_n] = [\vec{v}_1 \dots \vec{v}_n] \begin{bmatrix} \sigma_1^2 & & 0 \\ & \ddots & \\ 0 & & \sigma_n^2 \end{bmatrix}$
 $\Rightarrow [A^T A \vec{v}_1 \dots A^T A \vec{v}_n] = [\sigma_1^2 \vec{v}_1 \dots \sigma_n^2 \vec{v}_n]$
 $\Rightarrow A^T A \vec{v}_i = \sigma_i^2 \vec{v}_i, \dots, A^T A \vec{v}_n = \sigma_n^2 \vec{v}_n$
 $\vec{v}_i \sim \vec{v}_n$ 是相对于 $A^T A$ 的 eigenvector, 且 $A^T A$ 是 $n \times n$ symmetric eigenvector, 且 $A^T A$ 是 Hermitian. \Rightarrow 可正交对角化 $\Rightarrow n$ 个 orthonormal 的 eigenvector. 即 $\vec{v}_i \sim \vec{v}_n$

(2) $A = U \Sigma V^T$ (求 U)
 $\Rightarrow A V = U \Sigma$
 $\Rightarrow A [\vec{v}_1 \dots \vec{v}_n] = [u_1 \dots u_m] \begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_n \end{bmatrix}$
 $\Rightarrow A \vec{v}_i = \sigma_i u_i$
 $\Rightarrow \vec{u}_i = \frac{1}{\sigma_i} A \vec{v}_i, \quad i = 1, \dots, r$

(3) $A = \underbrace{[\vec{u}_1 \dots \vec{u}_r]}_{R(A)} \underbrace{\begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_r \end{bmatrix}}_{N(A^T)} \underbrace{[\vec{v}_1 \dots \vec{v}_n]^T}_{R(A^T)} \underbrace{[\vec{v}_{r+1} \dots \vec{v}_n]^T}_{N(A)}$ (秩 Σ)
 orthornormal basis

取 $\{\vec{u}_1, \dots, \vec{u}_r\}$ 为 $N(A^T)$ 之 一 组 orthornormal basis
 (可能导致 G.S. process)

Σ : 因 $A^T A$ 是正定, 其 eigenvalue ≥ 0 .
 据其 n 个 eigenvalue 找出 $\sigma_1 \sim \sigma_n$
 $|\vec{u}_i| \sim |\vec{v}_i|$ 即 Σ 对角线上
 singular value

I use five steps to get PCA.

step1: around center.

step2: $X.T@X$ to calculate sigma&V.

step3: sort eigenvalue big -> small.

step4: delete eigenvalue<0

step5: $U = (1/\text{sigma})X@V$

We can get:

X shape:(45045, 135)

V shape:(135, 135)

U shape:(45045, num_dim) # can set num_dim

```
# SVD = U@Sigma@V.T
def pca(X, num_dim=None):
    # step 1: around center
    X_mean = np.mean(X, axis=1).reshape(-1, 1)
    X_center = X - X_mean

    # step 2: X.T@X to calculate sigma&V
    eigenvalues, eigenvectors = np.linalg.eig(X_center.T @ X_center)
    print("V shape: ", eigenvectors.shape)

    # step 3: sort eigenvalue big 2 small
    sort_index = np.argsort(-eigenvalues)

    # step 4: delete eigenvalue<0
    if num_dim is None:
        for eigenvalue, i in zip(eigenvalues[sort_index],
                                np.arange(len(eigenvalues))):
            if eigenvalue <= 0:
                sort_index = sort_index[:i]
                break
    else:
        sort_index = sort_index[:num_dim]

    eigenvalues = eigenvalues[sort_index]

    # step 5: U = (1/sigma)X@V
    eigenvectors = X_center@eigenvectors[:, sort_index]
    eigenvectors_norm = np.linalg.norm(eigenvectors,axis=0)
    eigenvectors = eigenvectors/eigenvectors_norm
    print("U shape: ", eigenvectors.shape)

    return eigenvalues, eigenvectors, X_mean
```

Reconstruction and print result.

```
# reduce dim (projection)
Z = U.T@(X - X_mean)

# reconstruction
X_reconstruction = U@Z + X_mean
print_reconstruction(X, X_reconstruction, 10, H, W)
```

Use KNN to compare the performance.

```
def performance(X_test,y_test,Z_train,y_train,U,X_mean=None,k=3):
    if X_mean is None:
        X_mean=np.zeros((X_test.shape[0],1))

    # reduce dim (projection)
    Z_test=U.T@(X_test-X_mean)

    # k-nn
    predicted_y=np.zeros(Z_test.shape[1])
    for i in range(Z_test.shape[1]):
        distance=np.zeros(Z_train.shape[1])
        for j in range(Z_train.shape[1]):
            distance[j]=np.sum(np.square(Z_test[:,i]-Z_train[:,j]))
        sort_index=np.argsort(distance)
        nearest_neighbors=y_train[sort_index[:k]]
        unique, counts = np.unique(nearest_neighbors, return_counts=True)
        nearest_neighbors=[k for k,v in sorted(dict(zip(unique, counts)).items(),
                                                key=lambda item: -item[1])]

        predicted_y[i]=nearest_neighbors[0]

    acc=np.count_nonzero((y_test-predicted_y)==0)/len(y_test)
    return acc
```

4. LDA

Because of the singularity of the within-class scatter matrix. First I use pca to reduce dimension to 28, then use lda to reduce dimension to 27. It can get the best performance acc:60%.

The transformation matrix U is eigenvector of PCA @ eigenvector of LDA.

```
eigenvalues_pca, eigenvectors_pca, X_mean = pca(X, num_dim=num_dim)
X_pca = eigenvectors_pca.T@(X - X_mean)

eigenvalues_lda, eigenvectors_lda = lda(X_pca, y)

# Transform matrix
U = eigenvectors_pca@eigenvectors_lda
print('LDA U shape: {}'.format(U.shape))
```

```
V shape: (135, 135)
PCA U shape: (45045, 28)
LDA U shape: (45045, 27)
num_dim: 28, k: 9 ,acc: 60.00%
```

```

def lda(X,y,num_dim=None):
    N=X.shape[0]
    X_mean = np.mean(X, axis=1).reshape(-1, 1)

    classes_mean = np.zeros((N, 15)) # 15 classes
    for i in range(X.shape[1]):
        classes_mean[:, y[i]] += X[:, y[i]]
    classes_mean = classes_mean / 9

    # within-class scatter
    S_within = np.zeros((N, N))
    for i in range(X.shape[1]):
        d = X[:, y[i]].reshape(-1,1) - classes_mean[:, y[i]].reshape(-1,1)
        S_within += d @ d.T

    # between-class scatter
    S_between = np.zeros((N, N))
    for i in range(15):
        d = classes_mean[:, i].reshape(-1,1) - X_mean
        S_between += 9 * d @ d.T

    eigenvalues,eigenvectors=np.linalg.eig(np.linalg.inv(S_within)@S_between)
    sort_index=np.argsort(-eigenvalues)
    if num_dim is None:
        sort_index=sort_index[:-1] # reduce 1 dim
    else:
        sort_index=sort_index[:num_dim]

    eigenvalues=np.asarray(eigenvalues[sort_index].real,dtype='float')
    eigenvectors=np.asarray(eigenvectors[:,sort_index].real,dtype='float')

    return eigenvalues,eigenvectors

```

5. Kernel

Define rbf kernel & polynomial kernel.

```

def rbfkernel(X1, X2, gamma):
    return np.exp(-gamma*cdist(X1, X2, 'sqeuclidean'))

def polykernel(X):
    K = np.zeros(shape=(len(X), len(X)))
    for i in range(len(X)):
        for j in range(len(X)):
            k = 1 + np.dot(X[i].T,X[j]) ##  $K(i,j) = (1 + x(i).T \cdot x(j))^p$ 
            k = k**2
            K[i][j] = k
    return K

```

6. Kernel PCA

Through the kernel PCA, we perform a non-linear correspondence function to transform the data into a higher-dimensional space, and in this high-dimensional space, use standard PCA, and then project the data back to a lower-dimensional subspace.

```

def kpca(X, num_dim, kernel):
    # step0: kernel
    if kernel == 'RBF':
        gamma = 0.01
        X = rbfkernel(X, X, gamma)

    elif kernel == 'polynomial':
        X = polykernel(X)

    # step 1: around center
    X_mean = np.mean(X, axis=1).reshape(-1, 1)
    X_center = X - X_mean

    # step 2: X.T@X to caculate sigma&V
    eigenvalues, eigenvectors = np.linalg.eig(X_center.T @ X_center)
    print("V shape: ", eigenvectors.shape)

    # step 3: sort eigenvalue big 2 small
    sort_index = np.argsort(-eigenvalues)

    # step 4: delete eigenvalue<0
    if num_dim is None:
        for eigenvalue, i in zip(eigenvalues[sort_index],
                                np.arange(len(eigenvalues))):
            if eigenvalue <= 0:
                sort_index = sort_index[:i]
                break
    else:
        sort_index = sort_index[:num_dim]

    eigenvalues = eigenvalues[sort_index]

    # step 5: U = (1/sigma)X@V
    eigenvectors = X_center@eigenvectors[:, sort_index]
    eigenvectors_norm = np.linalg.norm(eigenvectors,axis=0)
    eigenvectors = eigenvectors/eigenvectors_norm
    print("U shape: ", eigenvectors.shape)

    return eigenvalues, eigenvectors, X_mean

```

7. Kernel LDA

Also through the kernel LDA, we perform a non-linear correspondence function to transform the data into a higher-dimensional space, and in this high-dimensional space, use standard LDA, and then project the data back to a lower-dimensional subspace.

```

def lda(X,y,num_dim=None):
    # step0: kernel
    if kernel == 'RBF':
        gamma = 0.01
        X = rbfkernel(X, X, gamma)

    elif kernel == 'polynomial':
        X = polykernel(X)

    N = X.shape[0]
    X_mean = np.mean(X, axis=1).reshape(-1, 1)

    classes_mean = np.zeros((N, 15)) # 15 classes
    for i in range(X.shape[1]):
        classes_mean[:, y[i]] += X[:, y[i]]
    classes_mean = classes_mean / 9

    # within-class scatter
    S_within = np.zeros((N, N))
    for i in range(X.shape[1]):
        d = X[:, y[i]].reshape(-1,1) - classes_mean[:, y[i]].reshape(-1,1)
        S_within += d @ d.T

    # between-class scatter
    S_between = np.zeros((N, N))
    for i in range(15):
        d = classes_mean[:, i].reshape(-1,1) - X_mean
        S_between += 9 * d @ d.T

    eigenvalues,eigenvectors=np.linalg.eig(np.linalg.inv(S_within)@S_between)
    sort_index=np.argsort(-eigenvalues)
    if num_dim is None:
        sort_index=sort_index[:-1] # reduce 1 dim
    else:
        sort_index=sort_index[:num_dim]

    eigenvalues=np.asarray(eigenvalues[sort_index].real,dtype='float')
    eigenvectors=np.asarray(eigenvectors[:,sort_index].real,dtype='float')

    return eigenvalues,eigenvectors

```


II. t-SNE

- Part1: Try to modify the code a little bit and make it back to symmetric SNE. You need to first understand how to implement t-SNE and find out the specific code piece to modify. You have to explain the difference between symmetric SNE and t-SNE in the report (e.g. point out the crowded problem of symmetric SNE).

- t-SNE & symmetric SNE is different in q_{ij} & **gradient**. t-SNE uses t- distribution. symmetric SNE uses Gaussian distribution.
- We can find $\|y_i - y_j\|^2 = np.add(np.add(num, sum_Y).T, sum_Y)$, so we modify num.

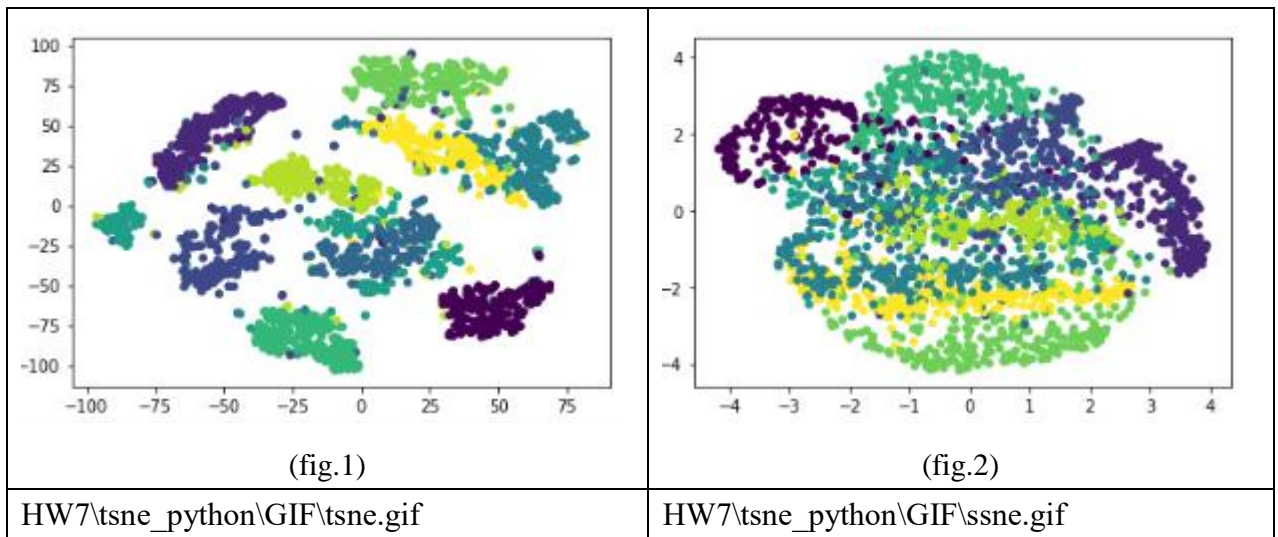
| t-SNE | symmetric SNE |
|--|---|
| $q_{ij} = \frac{(1 + \ y_i - y_j\ ^2)^{-1}}{\sum_{k \neq l} (1 + \ y_i - y_j\ ^2)^{-1}}$ | $q_{ij} = \frac{\exp(-\ y_i - y_j\ ^2)}{\sum_{k \neq l} \exp(-\ y_l - y_k\ ^2)}$ |
| <pre># Run iterations for iter in range(max_iter): # Compute pairwise affinities sum_Y = np.sum(np.square(Y), 1) num = -2. * np.dot(Y, Y.T) num = 1. / (1. + np.add(np.add(num, sum_Y).T, sum_Y)) num[range(n), range(n)] = 0. Q = num / np.sum(num) Q = np.maximum(Q, 1e-12)</pre> | <pre># Run iterations for iter in range(max_iter): # Compute pairwise affinities sum_Y = np.sum(np.square(Y), 1) num = -2. * np.dot(Y, Y.T) num = np.exp(-(np.add(np.add(num, sum_Y).T, sum_Y))) num[range(n), range(n)] = 0. Q = num / np.sum(num) Q = np.maximum(Q, 1e-12)</pre> |
| $\frac{\partial C}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \ y_i - y_j\ ^2)^{-1}$ | $\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{ij} - q_{ij})(y_i - y_j)$ |
| <pre># Compute gradient PQ = P - Q for i in range(n): dy[i, :] = np.sum(np.tile(PQ[:, i] * num[:, i], (no_dims, 1)).T * (Y[i, :] - Y), 0)</pre> | <pre># Compute gradient PQ = P - Q for i in range(n): dy[i, :] = np.dot(PQ[i, :], Y[i, :] - Y)</pre> |

- The **crowded problem** means that the clusters are gathered together and cannot be distinguished(fig.1).
- The distance relationship in high-dimensional data cannot be completely retained in low-dimensional space.

- Part2: Visualize the embedding of both t-SNE and symmetric SNE. Details of the visualization:
 - Project all your data onto 2D space and mark the data points into different colors respectively. The color of the data points depends on the label.
 - Use videos or GIF images to show the optimize procedure.

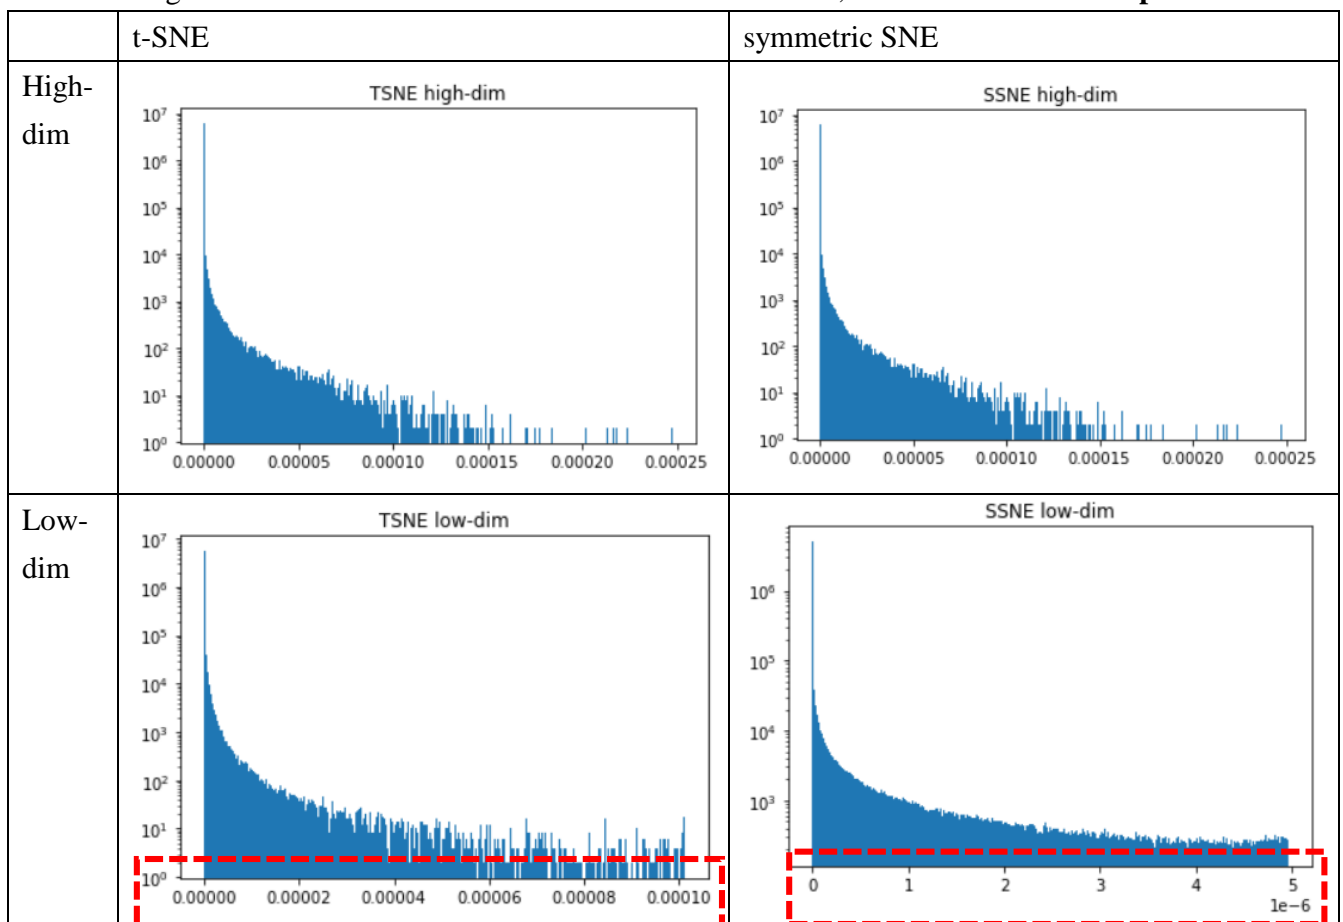
perplexity=20

| | |
|-------|---------------|
| t-SNE | symmetric SNE |
|-------|---------------|



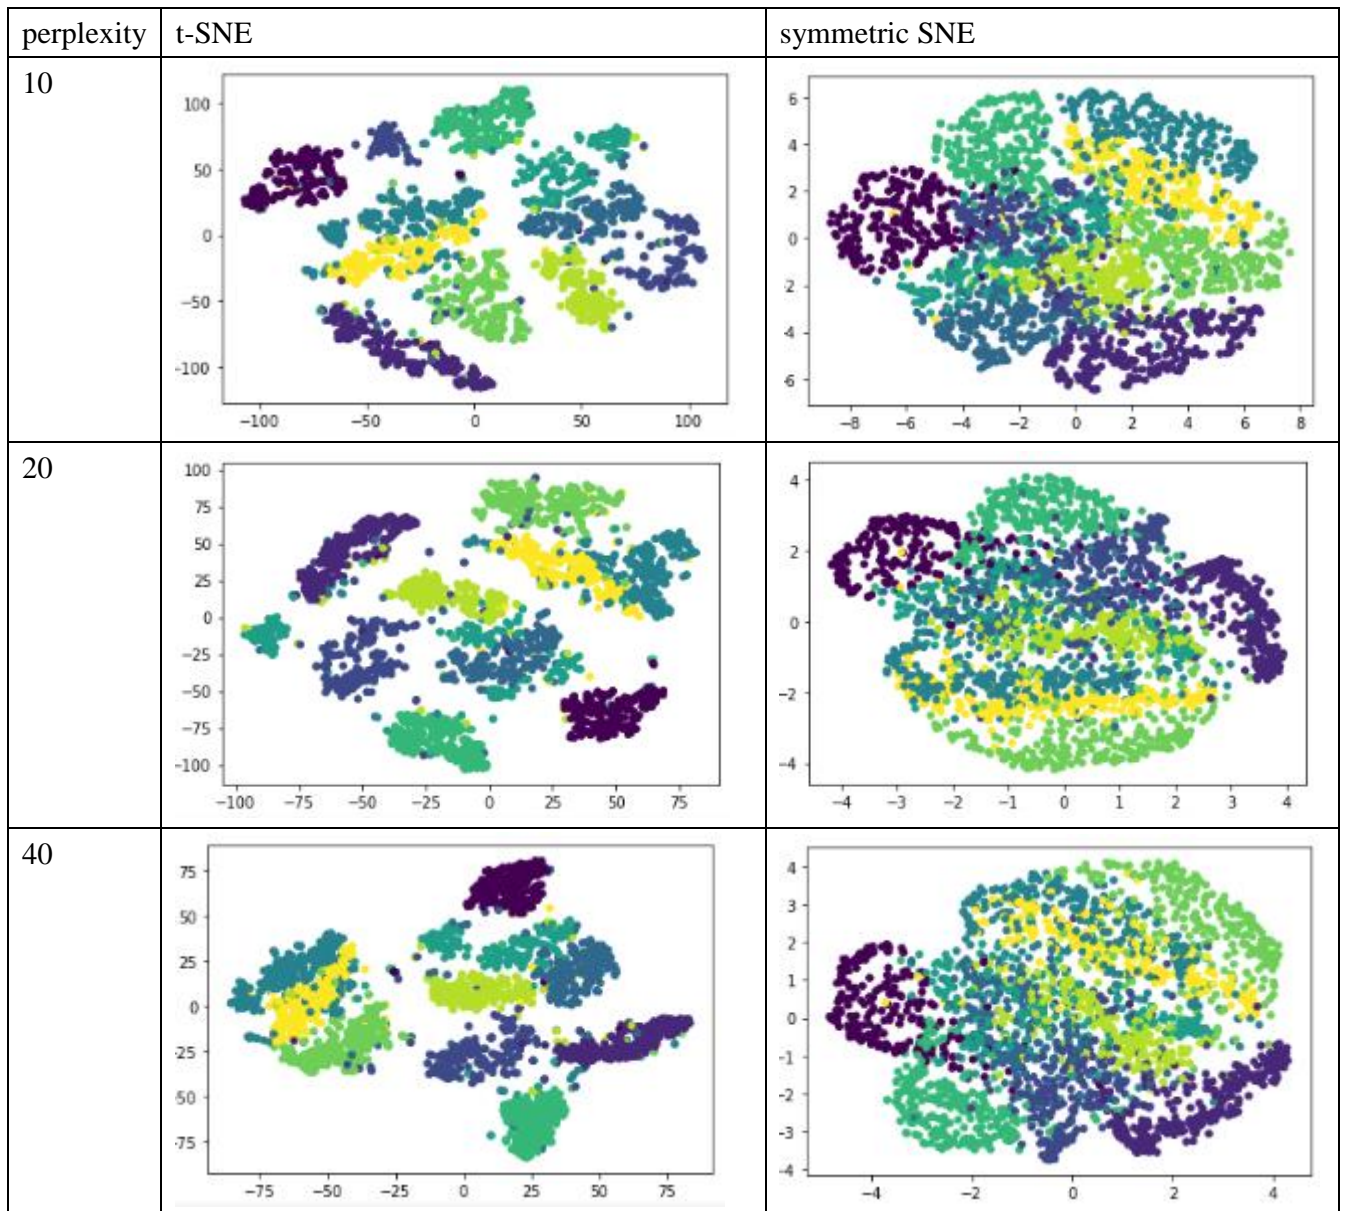
➤ Part3: Visualize the distribution of pairwise similarities in both high-dimensional space and low-dimensional space, based on both t-SNE and symmetric SNE.

- The range of ssne in low-dim is much smaller than that of tsne, so ssne has a **crowded problem**.



➤ Part4: Try to play with different perplexity values. Observe the change in visualization and explain it in the report.

- We observe a tendency towards clearer shapes as the perplexity value increases.



III. reference

- [1] https://www.youtube.com/watch?v=g-Hb26agBFg&ab_channel=Serrano.Academy
- [2] <https://kknews.cc/code/a6olrln.html>
- [3] <https://www.uj5u.com/houduan/294222.html>
- [4] <https://kknews.cc/code/o2gbm45.html>
- [5] <https://arbu00.blogspot.com/2017/02/7-kernel-pca.html>